



COMP250: Artificial Intelligence

7: Monte Carlo Tree Search

Learning outcomes

- ▶ Outcome 1
- ▶ Outcome 2
- ▶ Outcome 3

Heuristics for search



From session 2: Minimax search

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer is maximising then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue,  $v$ )
      if bestValue  $\geq 1$  then
        break
    return bestValue
  else if currentPlayer is minimising then
    bestValue =  $+\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MIN(bestValue,  $v$ )
      if bestValue  $\leq -1$  then
        break
  return bestValue
```

Minimax for larger games

Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states

Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully, e.g. chess has $\approx 10^{47}$ states

Heuristics

Heuristics

- ▶ A **heuristic** is an **approximate** solution to a problem, usually **quicker** to compute than a true solution

Depth limiting

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as $+1 / 0 / -1$

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as $+1 / 0 / -1$
- ▶ For nonterminal states at depth d , apply a heuristic evaluation instead of searching deeper

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as $+1 / 0 / -1$
- ▶ For nonterminal states at depth d , apply a heuristic evaluation instead of searching deeper
- ▶ Evaluation is a number between -1 and $+1$, estimating the probable outcome of the game

Move ordering

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search
- ▶ How do we know which moves are “best” and “worst”? Use a heuristic!