



COMP110: Principles of Computing **12: Machine Code**

Worksheets

- ▶ Worksheet 8 due **this Wednesday**
- ▶ Worksheet 9 due **in January**
- ▶ Summative submission of worksheets 1-9 due **in January** (check MyFalmouth for exact date!)
- ▶ Worksheet sign-off sessions **this week** and **4-6 January**

How programs are executed



Executing programs

- ▶ CPUs execute **machine code**

Executing programs

- ▶ CPUs execute **machine code**
- ▶ Programs must be **translated** into machine code for execution

Executing programs

- ▶ CPUs execute **machine code**
- ▶ Programs must be **translated** into machine code for execution
- ▶ There are three main ways of doing this:

Executing programs

- ▶ CPUs execute **machine code**
- ▶ Programs must be **translated** into machine code for execution
- ▶ There are three main ways of doing this:
 - ▶ An **interpreter** is an application which reads the program source code and executes it directly

Executing programs

- ▶ CPUs execute **machine code**
- ▶ Programs must be **translated** into machine code for execution
- ▶ There are three main ways of doing this:
 - ▶ An **interpreter** is an application which reads the program source code and executes it directly
 - ▶ An **ahead-of-time (AOT) compiler**, often just called a **compiler**, is an application which converts the program source code into executable machine code

Executing programs

- ▶ CPUs execute **machine code**
- ▶ Programs must be **translated** into machine code for execution
- ▶ There are three main ways of doing this:
 - ▶ An **interpreter** is an application which reads the program source code and executes it directly
 - ▶ An **ahead-of-time (AOT) compiler**, often just called a **compiler**, is an application which converts the program source code into executable machine code
 - ▶ A **just-in-time (JIT) compiler** is halfway between the two — it compiles the program on-the-fly at runtime

Examples

Interpreted:

- ▶ Python
- ▶ Lua
- ▶ Ruby
- ▶ Perl
- ▶ JavaScript
- ▶ GML
- ▶ Bespoke
scripting
languages

Examples

Interpreted:

- ▶ Python
- ▶ Lua
- ▶ Ruby
- ▶ Perl
- ▶ JavaScript
- ▶ GML
- ▶ Bespoke scripting languages

Compiled:

- ▶ C
- ▶ C++
- ▶ Swift
- ▶ Rust
- ▶ Go
- ▶ Blueprints
(via C++)

Examples

Interpreted:

- ▶ Python
- ▶ Lua
- ▶ Ruby
- ▶ Perl
- ▶ JavaScript
- ▶ GML
- ▶ Bespoke scripting languages

Compiled:

- ▶ C
- ▶ C++
- ▶ Swift
- ▶ Rust
- ▶ Go
- ▶ Blueprints (via C++)

JIT compiled:

- ▶ Java
- ▶ C#
- ▶ JavaScript

Examples

Interpreted:

- ▶ Python
- ▶ Lua
- ▶ Ruby
- ▶ Perl
- ▶ JavaScript
- ▶ GML
- ▶ Bespoke scripting languages

Compiled:

- ▶ C
- ▶ C++
- ▶ Swift
- ▶ Rust
- ▶ Go
- ▶ Blueprints (via C++)

JIT compiled:

- ▶ Java
- ▶ C#
- ▶ JavaScript

NB: technically any language could appear in any column here, but this is where they typically are

Interpreter vs compiler

- ▶ Run-time efficiency: compiler > interpreter

Interpreter vs compiler

- ▶ Run-time efficiency: compiler > interpreter
 - ▶ The compiler translates the program **in advance**, on the developer's machine

Interpreter vs compiler

- ▶ Run-time efficiency: compiler > interpreter
 - ▶ The compiler translates the program **in advance**, on the developer's machine
 - ▶ The interpreter translates the program **at runtime**, on the user's machine — this takes extra time

Interpreter vs compiler

- ▶ Portability: compiler < interpreter

Interpreter vs compiler

- ▶ Portability: compiler < interpreter
 - ▶ A compiled program can only run on the operating system and CPU architecture it was compiled for

Interpreter vs compiler

- ▶ Portability: compiler < interpreter
 - ▶ A compiled program can only run on the operating system and CPU architecture it was compiled for
 - ▶ An interpreted program can run on any machine, as long as a suitable interpreter is available

Interpreter vs compiler

- ▶ Ease of development: compiler < interpreter

Interpreter vs compiler

- ▶ Ease of development: compiler < interpreter
 - ▶ Writing an AOT or JIT compiler (especially a good one) is hard, and required in-depth knowledge of the target machine

Interpreter vs compiler

- ▶ Ease of development: compiler < interpreter
 - ▶ Writing an AOT or JIT compiler (especially a good one) is hard, and required in-depth knowledge of the target machine
 - ▶ Writing an interpreter is easy in comparison

Interpreter vs compiler

- ▶ Dynamic language features: compiler < interpreter

Interpreter vs compiler

- ▶ Dynamic language features: compiler < interpreter
 - ▶ The interpreter is already on the end user's machine, so programs can use it e.g. to dynamically generate and execute new code

Interpreter vs compiler

- ▶ Dynamic language features: compiler < interpreter
 - ▶ The interpreter is already on the end user's machine, so programs can use it e.g. to dynamically generate and execute new code
 - ▶ The AOT compiler is not generally on the end user's machine, so this is more difficult

Interpreter vs compiler

Interpreter vs compiler

- ▶ JIT compilers have similar pros/cons to interpreters

Interpreter vs compiler

- ▶ JIT compilers have similar pros/cons to interpreters
 - ▶ Runtime efficiency: JIT > interpreter (e.g. code inside a loop only needs to be translated once, then can be executed many times)

Interpreter vs compiler

- ▶ JIT compilers have similar pros/cons to interpreters
 - ▶ Runtime efficiency: JIT > interpreter (e.g. code inside a loop only needs to be translated once, then can be executed many times)
 - ▶ Ease of development: JIT < interpreter

Virtual machines

Virtual machines

- ▶ Many modern interpreters and JIT compilers translate programs into **bytecode**

Virtual machines

- ▶ Many modern interpreters and JIT compilers translate programs into **bytecode**
- ▶ Bytecode is essentially machine code for a **virtual machine (VM)**

Virtual machines

- ▶ Many modern interpreters and JIT compilers translate programs into **bytecode**
- ▶ Bytecode is essentially machine code for a **virtual machine (VM)**
- ▶ Translation from source code to bytecode can be done ahead of time

Virtual machines

- ▶ Many modern interpreters and JIT compilers translate programs into **bytecode**
- ▶ Bytecode is essentially machine code for a **virtual machine (VM)**
- ▶ Translation from source code to bytecode can be done ahead of time
- ▶ At runtime, translate the bytecode (by interpretation or JIT compilation) into machine code for the physical machine

Virtual machines

- ▶ Many modern interpreters and JIT compilers translate programs into **bytecode**
- ▶ Bytecode is essentially machine code for a **virtual machine (VM)**
- ▶ Translation from source code to bytecode can be done ahead of time
- ▶ At runtime, translate the bytecode (by interpretation or JIT compilation) into machine code for the physical machine
- ▶ E.g. a Java JAR file, a .NET executable, a Python .pyc or .pyo file all contain bytecode for their respective VMs

Assemblers

Assemblers

- ▶ **Assembly language** is designed to translate directly into machine code

Assemblers

- ▶ **Assembly language** is designed to translate directly into machine code
- ▶ An ahead-of-time compile for assembly language is called an **assembler**

Assemblers

- ▶ **Assembly language** is designed to translate directly into machine code
- ▶ An ahead-of-time compile for assembly language is called an **assembler**
- ▶ Generally much simpler than an AOT compiler for a higher-level language

The 6502 CPU architecture



MOS Technology 6502



MOS Technology 6502

- ▶ Designed by Chuck Peddle and team at **MOS Technology**



MOS Technology 6502



- ▶ Designed by Chuck Peddle and team at **MOS Technology**
 - ▶ **8-bit** CPU, 16-bit addressing

MOS Technology 6502



- ▶ Designed by Chuck Peddle and team at **MOS Technology**
- ▶ **8-bit CPU, 16-bit addressing**
- ▶ Clocked between **1MHz** and **3MHz**

MOS Technology 6502



- ▶ Designed by Chuck Peddle and team at **MOS Technology**
- ▶ **8-bit CPU**, 16-bit addressing
- ▶ Clocked between **1MHz** and **3MHz**
- ▶ First produced in **1975**

MOS Technology 6502



- ▶ Designed by Chuck Peddle and team at **MOS Technology**
- ▶ **8-bit CPU**, 16-bit addressing
- ▶ Clocked between **1MHz** and **3MHz**
- ▶ First produced in **1975**
- ▶ Still in production **today**

Uses of the 6502

Uses of the 6502



Uses of the 6502



Uses of the 6502



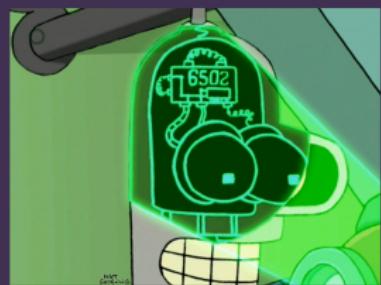
Uses of the 6502



Uses of the 6502



Uses of the 6502



Recap: hexadecimal notation

Recap: hexadecimal notation

- We usually write numbers in **decimal** i.e. **base 10**

Recap: hexadecimal notation

- ▶ We usually write numbers in **decimal** i.e. **base 10**
- ▶ Hexadecimal is **base 16**

Recap: hexadecimal notation

- ▶ We usually write numbers in **decimal** i.e. **base 10**
- ▶ Hexadecimal is **base 16**
- ▶ Uses extra digits:
 - ▶ A=10, B=11, ..., F=15

Recap: hexadecimal notation

- We usually write numbers in **decimal** i.e. **base 10**

- Hexadecimal is **base 16**

- Uses extra digits:

- A=10, B=11, ... , F=15

| Hex | Dec | Hex | Dec | Hex | Dec |
|-----|-----|-----|-----|-----|-----|
| 00 | 0 | 10 | 16 | F0 | 240 |
| 01 | 1 | 11 | 17 | F1 | 241 |
| : | : | : | : | : | : |
| 09 | 9 | 19 | 25 | F9 | 249 |
| 0A | 10 | 1A | 26 | FA | 250 |
| 0B | 11 | 1B | 27 | FB | 251 |
| 0C | 12 | 1C | 28 | FC | 252 |
| 0D | 13 | 1D | 29 | FD | 253 |
| 0E | 14 | 1E | 30 | FE | 254 |
| 0F | 15 | 1F | 31 | FF | 255 |

How a CPU works

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**
- ▶ An instruction is stored as an **opcode** followed by 0 or more **arguments**

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**
- ▶ An instruction is stored as an **opcode** followed by 0 or more **arguments**
- ▶ CPU has several **registers**, each storing a single value

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**
- ▶ An instruction is stored as an **opcode** followed by 0 or more **arguments**
- ▶ CPU has several **registers**, each storing a single value
 - ▶ Memory locations inside the CPU

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**
- ▶ An instruction is stored as an **opcode** followed by 0 or more **arguments**
- ▶ CPU has several **registers**, each storing a single value
 - ▶ Memory locations inside the CPU
 - ▶ Faster to access than main memory

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**
- ▶ An instruction is stored as an **opcode** followed by 0 or more **arguments**
- ▶ CPU has several **registers**, each storing a single value
 - ▶ Memory locations inside the CPU
 - ▶ Faster to access than main memory
- ▶ Instructions can read and write values in **registers** and in **memory**, and can perform **arithmetic and logical** operations on them

How a CPU works

- ▶ Executes a series of **instructions** stored in **memory**
- ▶ An instruction is stored as an **opcode** followed by 0 or more **arguments**
- ▶ CPU has several **registers**, each storing a single value
 - ▶ Memory locations inside the CPU
 - ▶ Faster to access than main memory
- ▶ Instructions can read and write values in **registers** and in **memory**, and can perform **arithmetic and logical** operations on them
- ▶ The **program counter (PC)** register stores the address of the next instruction to execute

Registers on the 6502

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255
- ▶ A = **accumulator**

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255
- ▶ A = **accumulator**
- ▶ X, Y = **index** registers

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255
- ▶ A = **accumulator**
- ▶ X, Y = **index** registers
- ▶ SP = **stack pointer** register

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255
- ▶ A = **accumulator**
- ▶ X, Y = **index** registers
- ▶ SP = **stack pointer** register
- ▶ PC = **program counter** register (16 bits)

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255
- ▶ A = **accumulator**
- ▶ X, Y = **index** registers
- ▶ SP = **stack pointer** register
- ▶ PC = **program counter** register (16 bits)
- ▶ **Status** register, composed of seven 1-bit **flags**

Registers on the 6502

- ▶ 6502 is an **8-bit** CPU, meaning each register stores **8 bits** of data
 - ▶ 1 byte
 - ▶ A number between 0 and 255
- ▶ A = **accumulator**
- ▶ X, Y = **index** registers
- ▶ SP = **stack pointer** register
- ▶ PC = **program counter** register (16 bits)
- ▶ **Status** register, composed of seven 1-bit **flags**
 - ▶ Includes information on the previous operation e.g. whether a calculation resulted in zero, negative, overflow...

Assembly language

Assembly language

- ▶ Translates **directly** to machine code

Assembly language

- ▶ Translates **directly** to machine code
- ▶ I.e. 1 line of assembly = 1 CPU instruction

Assembly language

- ▶ Translates **directly** to machine code
- ▶ I.e. 1 line of assembly = 1 CPU instruction
- ▶ An **assembler** translates assembly to machine code

Assembly language

- ▶ Translates **directly** to machine code
- ▶ I.e. 1 line of assembly = 1 CPU instruction
- ▶ An **assembler** translates assembly to machine code
 - ▶ I.e. an assembler is a “compiler” for assembly language

Assembly language

- ▶ Translates **directly** to machine code
- ▶ I.e. 1 line of assembly = 1 CPU instruction
- ▶ An **assembler** translates assembly to machine code
 - ▶ I.e. an assembler is a “compiler” for assembly language
- ▶ Each CPU architecture has its own instruction set therefore its own assembly language

Our first assembly program

Our first assembly program

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

Our first assembly program

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

Try it out! <http://skilldrick.github.io/easy6502/>

Our first assembly program

Our first assembly program

```
LDA #\$01
```

Our first assembly program

```
LDA #\$01
```

- Store the value 01 (hexadecimal) into register A

Our first assembly program

```
LDA #\$01
```

- ▶ Store the value 01 (hexadecimal) into register A
- ▶ **LDA** (“load accumulator”) stores a value in register A

Our first assembly program

```
LDA #\$01
```

- ▶ Store the value 01 (hexadecimal) into register A
- ▶ **LDA** (“load accumulator”) stores a value in register A
- ▶ # denotes a literal number (as opposed to a memory address)

Our first assembly program

```
LDA #\$01
```

- ▶ Store the value 01 (hexadecimal) into register A
- ▶ **LDA** (“load accumulator”) stores a value in register A
- ▶ # denotes a literal number (as opposed to a memory address)
- ▶ \$ denotes hexadecimal notation

Our first assembly program

Our first assembly program

STA \$0200

Our first assembly program

STA \$0200

- ▶ Write the value of register A into memory address 0200 (hex)

Our first assembly program

STA \$0200

- ▶ Write the value of register A into memory address 0200 (hex)
- ▶ **STA** ("store accumulator") copies the value of register A into main memory

Our first assembly program

STA \$0200

- ▶ Write the value of register A into memory address 0200 (hex)
- ▶ **STA** ("store accumulator") copies the value of register A into main memory
- ▶ Note that address is a **16-bit** number (2 bytes, 4 hex digits)

Our first assembly program

STA \$0200

- ▶ Write the value of register A into memory address 0200 (hex)
- ▶ **STA** ("store accumulator") copies the value of register A into main memory
- ▶ Note that address is a **16-bit** number (2 bytes, 4 hex digits)
- ▶ In this emulator the display is "memory mapped", with 1 byte per pixel, starting from address 0200
 - ▶ Real systems are usually more complicated than this!

Assembly to machine code

Assembly to machine code

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

Assembly to machine code

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

```
A9 01
8D 00 02
A9 05
8D 01 02
A9 08
8D 02 02
```

Assembly to machine code

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

```
A9 01
8D 00 02
A9 05
8D 01 02
A9 08
8D 02 02
```

Note that the 6502 is **little endian**

Assembly to machine code

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

```
A9 01
8D 00 02
A9 05
8D 01 02
A9 08
8D 02 02
```

Note that the 6502 is **little endian**

- ▶ In 16-bit values, the “low” byte comes before the “high” byte

Assembly to machine code

```
LDA #$01
STA $0200
LDA #$05
STA $0201
LDA #$08
STA $0202
```

```
A9 01
8D 00 02
A9 05
8D 01 02
A9 08
8D 02 02
```

Note that the 6502 is **little endian**

- ▶ In 16-bit values, the “low” byte comes before the “high” byte
- ▶ Intel x86 is also little endian

Looping

Looping

- ▶ PC normally **advances** to the next instruction

Looping

- ▶ PC normally **advances** to the next instruction
- ▶ Some instructions **modify** the PC

Looping

- ▶ PC normally **advances** to the next instruction
- ▶ Some instructions **modify** the PC
- ▶ E.g. **JMP** (jump) sets the PC to the specified address

Looping

- ▶ PC normally **advances** to the next instruction
- ▶ Some instructions **modify** the PC
- ▶ E.g. **JMP** (jump) sets the PC to the specified address

INC \$0200 ; add 1 to the value at address 0200

JMP \$0600 ; jump back to beginning of program

Looping

- ▶ PC normally **advances** to the next instruction
- ▶ Some instructions **modify** the PC
- ▶ E.g. **JMP** (jump) sets the PC to the specified address

```
INC $0200 ; add 1 to the value at address 0200
```

```
JMP $0600 ; jump back to beginning of program
```

- ▶ In this emulator the program always starts at address 0600
 - ▶ This may **not** be the case on other 6502-based systems!

Labels

Labels

- ▶ Don't use explicit jump locations in your code, it's not maintainable!

Labels

- ▶ Don't use explicit jump locations in your code, it's not maintainable!
- ▶ Can add a **label** to a line of code, by giving a name followed by a colon

Labels

- ▶ Don't use explicit jump locations in your code, it's not maintainable!
- ▶ Can add a **label** to a line of code, by giving a name followed by a colon
- ▶ Labels can then be used in instructions

Labels

- ▶ Don't use explicit jump locations in your code, it's not maintainable!
- ▶ Can add a **label** to a line of code, by giving a name followed by a colon
- ▶ Labels can then be used in instructions

```
start:  
    INC $0200  
    JMP start
```

Labels

- ▶ Don't use explicit jump locations in your code, it's not maintainable!
- ▶ Can add a **label** to a line of code, by giving a name followed by a colon
- ▶ Labels can then be used in instructions

```
start:  
  INC $0200  
  JMP start
```

- ▶ `start` is essentially a constant with value `$0600`

Labels

- ▶ Don't use explicit jump locations in your code, it's not maintainable!
- ▶ Can add a **label** to a line of code, by giving a name followed by a colon
- ▶ Labels can then be used in instructions

```
start:  
    INC $0200  
    JMP start
```

- ▶ `start` is essentially a constant with value `$0600`
- ▶ The assembled code is exactly the same as for the previous slide

Conditional branching

Conditional branching

```
LDX #$08      ; set X=8
decrement:
DEX          ; subtract 1 from X
STX $0200    ; store X in top left pixel
CPX #$03    ; compare X to 3
BNE decrement ; if not equal, jump
STX $0201    ; store X in next pixel
BRK          ; halt execution
```

Conditional branching

```
LDX #$08      ; set X=8
decrement:
DEX          ; subtract 1 from X
STX $0200    ; store X in top left pixel
CPX #$03    ; compare X to 3
BNE decrement ; if not equal, jump
STX $0201    ; store X in next pixel
BRK          ; halt execution
```

$X = 8$

do

$X = X - 1$

$\text{memory}[0200] = X$

while $X \neq 3$

$\text{memory}[0201] = X$

Conditional branching

Conditional branching

- ▶ Assembly language does not have **structured programming** constructs such as if/else, switch/case, for, while, etc.

Conditional branching

- ▶ Assembly language does not have **structured programming** constructs such as if/else, switch/case, for, while, etc.
- ▶ However all of these can be implemented using branch instructions

Conditional branching

- ▶ Assembly language does not have **structured programming** constructs such as if/else, switch/case, for, while, etc.
- ▶ However all of these can be implemented using branch instructions
- ▶ ... which is exactly how compilers implement them

Conditionals

Conditionals

- Branching allows us to implement **if statements**

Conditionals

- Branching allows us to implement **if statements**

```
CPX #$01      ; compare X to 1
BMI else      ; if X < 1, jump
DEY           ; Y = Y - 1
JMP end       ; skip over else block
else:
    INY          ; Y = Y + 1
end:
```

Conditionals

- Branching allows us to implement **if statements**

```
CPX #$01      ; compare X to 1
BMI else      ; if X < 1, jump
DEY           ; Y = Y - 1
JMP end       ; skip over else block
else:
    INY          ; Y = Y + 1
end:
```

```
if X ≥ 1 then
    Y = Y - 1
else
    Y = Y + 1
end if
```

Subroutines

Subroutines

- ▶ **JSR** (jump to subroutine) works like **JMP**, but stores the current PC

Subroutines

- ▶ **JSR** (jump to subroutine) works like **JMP**, but stores the current PC
- ▶ **RTS** (return from subroutine) jumps back to the instruction after the **JSR**

Subroutines

- ▶ `JSR` (jump to subroutine) works like `JMP`, but stores the current PC
- ▶ `RTS` (return from subroutine) jumps back to the instruction after the `JSR`
- ▶ These are used to implement **function calls**

Subroutines

- ▶ **JSR** (jump to subroutine) works like **JMP**, but stores the current PC
- ▶ **RTS** (return from subroutine) jumps back to the instruction after the **JSR**
- ▶ These are used to implement **function calls**
- ▶ Return addresses are stored on a **stack**

Addressing modes

Addressing modes

- ▶ Immediate: `LDA #$42`

Addressing modes

- ▶ Immediate: **LDA** #\$42
 - ▶ Load the literal value 42 (hex) into register A

Addressing modes

- ▶ Immediate: `LDA #$42`
 - ▶ Load the literal value 42 (hex) into register A
- ▶ Absolute: `LDA $42`

Addressing modes

- ▶ Immediate: `LDA #$42`
 - ▶ Load the literal value 42 (hex) into register A
- ▶ Absolute: `LDA $42`
 - ▶ Load the value stored at memory address 42 (hex) into register A

Addressing modes

- ▶ Immediate: `LDA #$42`
 - ▶ Load the literal value 42 (hex) into register A
- ▶ Absolute: `LDA $42`
 - ▶ Load the value stored at memory address 42 (hex) into register A
- ▶ That `#` makes a big difference!

Addressing modes

- ▶ Immediate: `LDA #$42`
 - ▶ Load the literal value 42 (hex) into register A
- ▶ Absolute: `LDA $42`
 - ▶ Load the value stored at memory address 42 (hex) into register A
- ▶ That `#` makes a big difference!
- ▶ Note that these actually assemble to **different** CPU instructions

Indexed addressing

```
LDA $0200,X
```

Indexed addressing

```
LDA $0200,X
```

- ▶ Look up the value stored at memory address
 $0200 + (\text{value of X register})$
and store it in A

Indexed addressing

```
LDA $0200,X
```

- ▶ Look up the value stored at memory address
 $0200 + (\text{value of X register})$
and store it in A
- ▶ Can also do LDA \$0200,Y

Indexed addressing

LDA \$0200,X

- ▶ Look up the value stored at memory address
 $0200 + (\text{value of X register})$
and store it in A
- ▶ Can also do **LDA** \$0200,Y
- ▶ ... but **only** X and Y registers can be used for indexed addressing

Indexed addressing

```
LDX #0          ; X=0
loop:
TXA            ; A=X
STA $0200,X    ; store A to 0200+X
INX            ; X++
JMP loop      ; loop forever
```

Indexed addressing

```
LDX #0          ; X=0
loop:
TXA            ; A=X
STA $0200,X    ; store A to 0200+X
INX            ; X++
JMP loop      ; loop forever
```

- ▶ Why does it stop $\frac{1}{4}$ of the way down?

Indexed addressing

```
LDX #0          ; X=0
loop:
TXA            ; A=X
STA $0200,X    ; store A to 0200+X
INX            ; X++
JMP loop      ; loop forever
```

- ▶ Why does it stop $\frac{1}{4}$ of the way down?
- ▶ Hint: it stops after filling 256 pixels...

The x86 CPU Architecture



Intel x86

Intel x86

- ▶ Introduced in 1978 with the Intel 8086 (16-bit)

Intel x86

- ▶ Introduced in 1978 with the Intel 8086 (16-bit)
- ▶ Extended to 32-bit in 1985 with the Intel 80306

Intel x86

- ▶ Introduced in 1978 with the Intel 8086 (16-bit)
- ▶ Extended to 32-bit in 1985 with the Intel 80386
- ▶ Extended to 64-bit (x86-64 or x64) in 2003 with the AMD Opteron

Uses of Intel x86

Uses of Intel x86



Uses of Intel x86



Uses of Intel x86



Uses of Intel x86



Uses of Intel x86



Uses of Intel x86



Backwards compatibility

Backwards compatibility

- ▶ New CPUs have built upon the x86 instruction set over the years

Backwards compatibility

- ▶ New CPUs have built upon the x86 instruction set over the years
- ▶ x86-64 includes 64, 32 and 16-bit instructions

Backwards compatibility

- ▶ New CPUs have built upon the x86 instruction set over the years
- ▶ x86-64 includes 64, 32 and 16-bit instructions
- ▶ Technically, a PC from 2020 could run the same OS and software as the IBM 5150

Features of modern x86-64

Features of modern x86-64

- ▶ x87 floating point unit

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data
 - ▶ E.g. one instruction to multiply 4 pairs of numbers in parallel

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data
 - ▶ E.g. one instruction to multiply 4 pairs of numbers in parallel
- ▶ Security features

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data
 - ▶ E.g. one instruction to multiply 4 pairs of numbers in parallel
- ▶ Security features
 - ▶ AES encryption and decryption

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data
 - ▶ E.g. one instruction to multiply 4 pairs of numbers in parallel
- ▶ Security features
 - ▶ AES encryption and decryption
 - ▶ Secure random number generation

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data
 - ▶ E.g. one instruction to multiply 4 pairs of numbers in parallel
- ▶ Security features
 - ▶ AES encryption and decryption
 - ▶ Secure random number generation
- ▶ Multiple cores and hyperthreading

Features of modern x86-64

- ▶ x87 floating point unit
 - ▶ Hardware support for IEEE754 floating point
- ▶ SIMD instructions
 - ▶ SIMD = Single Instruction, Multiple Data
 - ▶ E.g. one instruction to multiply 4 pairs of numbers in parallel
- ▶ Security features
 - ▶ AES encryption and decryption
 - ▶ Secure random number generation
- ▶ Multiple cores and hyperthreading
 - ▶ CPU can do more than one thing at once

RISC vs CISC





CISC

CISC

- x86 is what's known as a **CISC architecture**

CISC

- ▶ x86 is what's known as a **CISC architecture**
- ▶ CISC = Complex Instruction Set Computer

CISC

- ▶ x86 is what's known as a **CISC architecture**
- ▶ CISC = Complex Instruction Set Computer
- ▶ x86-64 has **thousands** of instructions

RISC

RISC

- ▶ RISC = Reduced Instruction Set Computer

RISC

- ▶ RISC = Reduced Instruction Set Computer
- ▶ E.g. ARM (Advanced RISC Machine) architecture

RISC

- ▶ RISC = Reduced Instruction Set Computer
- ▶ E.g. ARM (Advanced RISC Machine) architecture
- ▶ Widely used in mobile (Apple and Android)

RISC

- ▶ RISC = Reduced Instruction Set Computer
- ▶ E.g. ARM (Advanced RISC Machine) architecture
- ▶ Widely used in mobile (Apple and Android)
- ▶ **Tens or hundreds** of instructions

RISC vs CISC

RISC vs CISC

- ▶ CISC generally leads to **shorter programs** — CISC does more per instruction

RISC vs CISC

- ▶ CISC generally leads to **shorter programs** — CISC does more per instruction
- ▶ RISC generally leads to **simpler hardware** — instructions are simpler

RISC vs CISC

- ▶ CISC generally leads to **shorter programs** — CISC does more per instruction
- ▶ RISC generally leads to **simpler hardware** — instructions are simpler
- ▶ This also leads to better efficiency in terms of **power** and **heat**

RISC vs CISC

- ▶ CISC generally leads to **shorter programs** — CISC does more per instruction
- ▶ RISC generally leads to **simpler hardware** — instructions are simpler
- ▶ This also leads to better efficiency in terms of **power** and **heat**
- ▶ Assembly programmers and compiler developers need to worry about the difference — it is abstracted away from the rest of us

RISC vs CISC

- ▶ CISC generally leads to **shorter programs** — CISC does more per instruction
- ▶ RISC generally leads to **simpler hardware** — instructions are simpler
- ▶ This also leads to better efficiency in terms of **power** and **heat**
- ▶ Assembly programmers and compiler developers need to worry about the difference — it is abstracted away from the rest of us
- ▶ Nowadays the main reason to use CISC is **backwards compatibility**

RISC vs CISC

- ▶ CISC generally leads to **shorter programs** — CISC does more per instruction
- ▶ RISC generally leads to **simpler hardware** — instructions are simpler
- ▶ This also leads to better efficiency in terms of **power** and **heat**
- ▶ Assembly programmers and compiler developers need to worry about the difference — it is abstracted away from the rest of us
- ▶ Nowadays the main reason to use CISC is **backwards compatibility**
- ▶ Apple are moving to RISC for Mac computers — will PCs ever follow suit?

Workshop



Workshop

- ▶ Work through the Easy 6502 article/tutorial linked on LearningSpace
- ▶ Start work on Worksheet 9 (TIS-100)