



FALMOUTH  
UNIVERSITY



COMP250: Artificial Intelligence

## 3: Planning

# Game theory



# Game theory

# Game theory

- ▶ A branch of mathematics studying **decision making**

# Game theory

- ▶ A branch of mathematics studying **decision making**
- ▶ A **game** is a system where one or more **players** choose **actions**; the combination of these choices lead to each agent receiving a **payoff**

# Game theory

- ▶ A branch of mathematics studying **decision making**
- ▶ A **game** is a system where one or more **players** choose **actions**; the combination of these choices lead to each agent receiving a **payoff**
- ▶ Important applications in economics, ecology and social sciences as well as AI

# The Prisoner's Student's Dilemma

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other



# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled
- ▶ If **Bob betrays Alice**, he receives an A whilst she gets expelled

# The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled
- ▶ If **Bob betrays Alice**, he receives an A whilst she gets expelled
- ▶ If **both betray each other**, both get an F

# Payoff matrix

# Payoff matrix

	A silent	A betray
B silent	A: 50 B: 50	A: 70 B: -100
B betray	A: -100 B: 70	A: 0 B: 0



# Alice's thought process

# Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent

# Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
- ▶ However if Bob stays silent, I can get a better mark by betraying him

# Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
- ▶ However if Bob stays silent, I can get a better mark by betraying him
- ▶ And if Bob betrays me, I should betray him to avoid getting expelled

# Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
- ▶ However if Bob stays silent, I can get a better mark by betraying him
- ▶ And if Bob betrays me, I should betray him to avoid getting expelled
- ▶ Therefore the rational choice is to betray

# Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
  - ▶ However if Bob stays silent, I can get a better mark by betraying him
  - ▶ And if Bob betrays me, I should betray him to avoid getting expelled
  - ▶ Therefore the rational choice is to betray
- ... and Bob's thought process is the same!

# Nash equilibrium

# Nash equilibrium

- ▶ Consider the situation where both have chosen to betray



# Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch

# Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch
- ▶ Such a situation is called a **Nash equilibrium**

# Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch
- ▶ Such a situation is called a **Nash equilibrium**
- ▶ If all players are **rational** (in the sense of wanting to maximising payoff), they should converge upon a Nash equilibrium

# Does every game have a Nash equilibrium?

# Does every game have a Nash equilibrium?

	A rock	A paper	A scissors
B rock	A: 0 B: 0	A: +1 B: -1	A: -1 B: +1
B paper	A: -1 B: +1	A: 0 B: 0	A: +1 B: -1
B scissors	A: +1 B: -1	A: -1 B: +1	A: 0 B: 0

# Nash equilibrium for Rock-Paper-Scissors

# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**

# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors



# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop

# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
  - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...

# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
  - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...
- ▶ The optimum strategy is to be **unpredictable**

# Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
  - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...
- ▶ The optimum strategy is to be **unpredictable**
- ▶ Choose rock with probability  $\frac{1}{3}$ , paper with probability  $\frac{1}{3}$ , scissors with probability  $\frac{1}{3}$

# Mixed strategies

# Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random

# Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random
- ▶ In contrast to a **pure** or **deterministic strategy**, which always chooses the same action

# Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random
- ▶ In contrast to a **pure** or **deterministic strategy**, which always chooses the same action
- ▶ If we allow mixed strategies, **every game has at least one Nash equilibrium**



Guess  $\frac{2}{3}$  of the average

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to  $\frac{2}{3}$  of the mean of all guesses

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to  $\frac{2}{3}$  of the mean of all guesses
- ▶ Example:

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to  $\frac{2}{3}$  of the mean of all guesses
- ▶ Example:
  - ▶ If the guesses are 30, 40 and 80...

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to  $\frac{2}{3}$  of the mean of all guesses
- ▶ Example:
  - ▶ If the guesses are 30, 40 and 80...
  - ▶ ... then the mean is  $\frac{30+40+80}{3} = 50...$

# Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to  $\frac{2}{3}$  of the mean of all guesses
- ▶ Example:
  - ▶ If the guesses are 30, 40 and 80...
  - ▶ ... then the mean is  $\frac{30+40+80}{3} = 50...$
  - ▶ ... so the winning guess is 30, as this is closest to  $\frac{2}{3} \times 50 = 33.333$

# The rational guess



# The rational guess

- ▶ The average can't possibly be greater than 100

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629
- ▶ ... and so on ad infinitum

# The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629
- ▶ ... and so on ad infinitum
- ▶ So the only **rational** guess is 0, as every rational player should guess 0 and  $\frac{2}{3}$  of 0 is 0



# Rationality

# Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers

# Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers
- ▶ However it's important to remember that **humans aren't always rational**

# Planning



# Planning

# Planning

- ▶ An **agent** in an **environment**

# Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**

# Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state



# Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**

# Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**
- ▶ Problem: find a sequence of actions that leads to the goal

# STRIPS planning

# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver

# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true

# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:

# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
  - ▶ The **initial state** (a set of predicates which are true)

# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
  - ▶ The **initial state** (a set of predicates which are true)
  - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)



# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
  - ▶ The **initial state** (a set of predicates which are true)
  - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
  - ▶ The set of **actions**, each specifying:

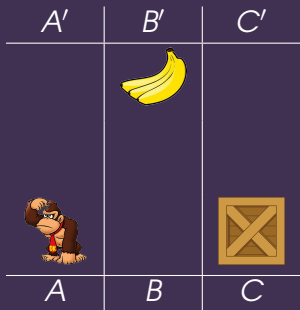
# STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
  - ▶ The **initial state** (a set of predicates which are true)
  - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
  - ▶ The set of **actions**, each specifying:
    - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)

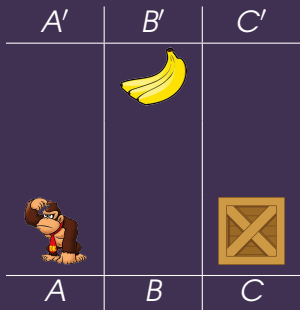
# STRIPS planning

- ▶ **Stanford Research Institute Problem Solver**
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
  - ▶ The **initial state** (a set of predicates which are true)
  - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
  - ▶ The set of **actions**, each specifying:
    - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)
    - ▶ Postconditions (specifying what predicates are made true or false by this action)

# STRIPS example




# STRIPS example



Initial state:

At (A) ,  
BoxAt (C) ,  
BananasAt (B' )

# STRIPS example




Initial state:

At (A) ,  
BoxAt (C) ,  
BananasAt (B' )

Goal:

HasBananas

# STRIPS example — Actions



Move(x, y)

Pre: At(x)

Post: !At(x), At(y)

ClimbUp(x)

Pre: At(x), BoxAt(x)

Post: !At(x), At(x')

ClimbDown(x')

Pre: At(x'), BoxAt(x)

Post: !At(x'), At(x)

PushBox(x, y)

Pre: At(x), BoxAt(x)


Post: !At(x), At(y),  
!BoxAt(x), BoxAt(y)

TakeBananas(x)

Pre: At(x), BananasAt(x)


Post: !BananasAt(x), HasBananas

# STRIPS example — Solution







# STRIPS example — Solution




# STRIPS example — Solution



# STRIPS example — Solution



# STRIPS example — Solution



# Finding the solution

# Finding the solution

- For a given state, we can construct a list of all **valid actions** based on their **preconditions**

# Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**

# Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**



# Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
  - ▶ Nodes: environment states

# Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
  - ▶ Nodes: environment states
  - ▶ Edges: actions

# Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
  - ▶ Nodes: environment states
  - ▶ Edges: actions
- ▶ We can then **search** this tree to find a goal state

# Searching for the solution

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?



# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
  - ▶ Depth-first search

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
  - ▶ Depth-first search
  - ▶ Breadth-first search

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
  - ▶ Depth-first search
  - ▶ Breadth-first search
  - ▶ Dijkstra's algorithm

# Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
  - ▶ Depth-first search
  - ▶ Breadth-first search
  - ▶ Dijkstra's algorithm
  - ▶ A\* (if we have a suitable **heuristic**)