



COMP250: Artificial Intelligence

3: Game Theory and Planning

Game theory

Game theory

- ▶ A branch of mathematics studying **decision making**
- ▶ A **game** is a system where one or more **players** choose **actions**; the combination of these choices lead to each agent receiving a **payoff**
- ▶ Important applications in economics, ecology and social sciences as well as AI

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled
- ▶ If **Bob betrays Alice**, he receives an A whilst she gets expelled
- ▶ If **both betray each other**, both get an F

Payoff matrix

	A silent	A betray
B silent	A: 50 B: 50	A: 70 B: -100
B betray	A: -100 B: 70	A: 0 B: 0

Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch
- ▶ Such a situation is called a **Nash equilibrium**
- ▶ If all players are **rational** (in the sense of wanting to maximising payoff), they should converge upon a Nash equilibrium

Does every game have a Nash equilibrium?

	A rock	A paper	A scissors
B rock	A: 0 B: 0	A: +1 B: -1	A: -1 B: +1
B paper	A: -1 B: +1	A: 0 B: 0	A: +1 B: -1
B scissors	A: +1 B: -1	A: -1 B: +1	A: 0 B: 0

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
 - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...
- ▶ The optimum strategy is to be **unpredictable**
- ▶ Choose rock with probability $\frac{1}{3}$, paper with probability $\frac{1}{3}$, scissors with probability $\frac{1}{3}$

Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random
- ▶ In contrast to a **pure** or **deterministic strategy**, which always chooses the same action
- ▶ If we allow mixed strategies, **every game has at least one Nash equilibrium**

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to $\frac{2}{3}$ of the mean of all guesses
- ▶ Example:
 - ▶ If the guesses are 30, 40 and 80...
 - ▶ ... then the mean is $\frac{30+40+80}{3} = 50...$
 - ▶ ... so the winning guess is 30, as this is closest to $\frac{2}{3} \times 50 = 33.333$

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629
- ▶ ... and so on ad infinitum
- ▶ So the only **rational** guess is 0, as every rational player should guess 0 and $\frac{2}{3}$ of 0 is 0

Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers
- ▶ However it's important to remember that **humans aren't always rational**
- ▶ E.g. in guess $\frac{2}{3}$, if enough players deviate from the rational guess of 0, then 0 is no longer the best guess!

Planning


Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**
- ▶ Problem: find a sequence of actions that leads to the goal

STRIPS planning

- ▶ **S**tanford **R**esearch Institute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:
 - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)
 - ▶ Postconditions (specifying what predicates are made true or false by this action)

STRIPS example




Initial state:

At (A) ,
BoxAt (C) ,
BananasAt (B')

Goal:

HasBananas

STRIPS example — Actions



`Move(x, y)`

Pre: `At(x)`

Post: `!At(x), At(y)`

`ClimbUp(x)`

Pre: `At(x), BoxAt(x)`

Post: `!At(x), At(x')`

`ClimbDown(x')`

Pre: `At(x'), BoxAt(x)`

Post: `!At(x'), At(x)`

`PushBox(x, y)`

Pre: `At(x), BoxAt(x)`


Post: `!At(x), At(y),
!BoxAt(x), BoxAt(y)`

`TakeBananas(x)`

Pre: `At(x), BananasAt(x)`

Post: `!BananasAt(x), HasBananas`


STRIPS example — Solution



State:

At (A) ,
BoxAt (C) ,
BananasAt (B')

STRIPS example — Solution




Action:

Move (A, C)

State:

At (C) ,
BoxAt (C) ,
BananasAt (B')

STRIPS example — Solution




Action:

PushBox (C, B)

State:

At (B) ,
BoxAt (B) ,
BananasAt (B')

STRIPS example — Solution

A'	B'	C'
		
A	B	C


Action:

ClimbUp (B)

State:

At (B') ,
BoxAt (B) ,
BananasAt (B')

STRIPS example — Solution



Action:

TakeBananas (B')

State:

At (B'),
BoxAt (B),
HasBananas

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
 - ▶ Nodes: environment states
 - ▶ Edges: actions
- ▶ We can then **search** this graph to find the optimum path from the current state to a goal state

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
 - ▶ Depth-first search
 - ▶ Breadth-first search
 - ▶ Dijkstra's algorithm (same as BFS unless actions have non-uniform **cost**)
 - ▶ A* (if we have a suitable **heuristic**)

GOAP

GOAP

- ▶ **G**oal **O**riented **A**ction **P**lanning
- ▶ Originally developed for F.E.A.R. (2005), since used in several games
- ▶ A modified version of STRIPS specifically for real-time planning in video games

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**
 - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
 - ▶ Like in STRIPS — actions have preconditions and postconditions
 - ▶ Unlike STRIPS, each action also has a **cost**

Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system
- ▶ E.g. this was used by the F.E.A.R. team to quickly add new enemy types

Action sets

Soldier

- 1 AI/Actions/Attack
- 2 AI/Actions/AttackCrouch
- 3 AI/Actions/SuppressionFire
- 4 AI/Actions/SuppressionFireFromCover
- 5 AI/Actions/FlushOutWithGrenade
- 6 AI/Actions/AttackFromCover
- 7 AI/Actions/BlindFireFromCover
- 8 AI/Actions/AttackGrenadeFromCover
- 9 AI/Actions/AttackFromView
- 10 AI/Actions/DrawWeapon
- 11 AI/Actions/HolsterWeapon
- 12 AI/Actions/ReloadCrouch
- 13 AI/Actions/ReloadCovered
- 14 AI/Actions/InspectDisturbance
- 15 AI/Actions/LookAtDisturbance
- 16 AI/Actions/SurveyArea
- 17 AI/Actions/DodgeRoll
- 18 AI/Actions/DodgeShuffle
- 19 AI/Actions/DodgeCovered
- 20 AI/Actions/Uncover
- 21 AI/Actions/AttackMelee

Assassin

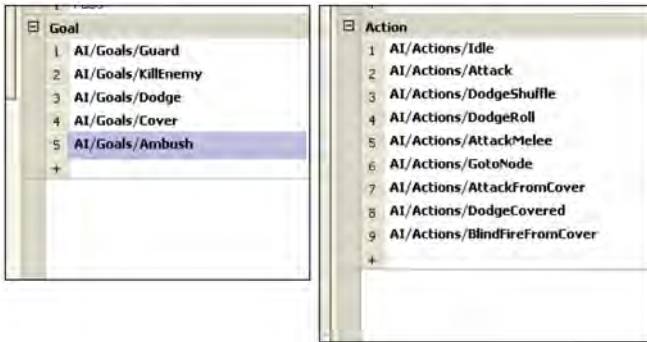
- 1 AI/Actions/Attack
- 2 AI/Actions/InspectDisturbance
- 3 AI/Actions/LookAtDisturbance
- 4 AI/Actions/SurveyArea
- 5 AI/Actions/AttackMeleeUncloaked
- 6 AI/Actions/TraverseBlockedDoor
- 7 AI/Actions/UseSmartObjectNodeMounted
- 8 AI/Actions/MountNodeUncloaked
- 9 AI/Actions/DismountNodeUncloaked
- 10 AI/Actions/TraverseLinkUncloaked
- 11 AI/Actions/AttackFromAmbush
- 12 AI/Actions/DodgeRollParanoid
- 13 AI/Actions/AttackLungeUncloaked
- 14 AI/Actions/LopeToTargetUncloaked

Rat

- 1 AI/Actions/Animate
- 2 AI/Actions/Idle
- 3 AI/Actions/GotoNode
- 4 AI/Actions/UseSmartObjectNode

Layering

- ▶ Goal set allows different behaviours with different priorities to be **layered**
- ▶ E.g. enemy AI in F.E.A.R.:



Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation
- ▶ Most implementations also allow for **programmatic** preconditions (e.g. calling the pathfinding system to check availability of a path)

Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist
- ▶ Not built into Unity or Unreal, but asset store packages are available

Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use Dijkstra or A* to find the lowest cost path to the goal
- ▶ Plan is a **queue** of actions that the agent then executes
- ▶ If the plan is interrupted or fails then the agent can **replan**

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**
- ▶ GOAP can be classified as **computational intelligence**