



COMP250: Artificial Intelligence

3: Game Theory and Planning



Game theory



Game theory

Game theory

- ▶ A branch of mathematics studying **decision making**

Game theory

- ▶ A branch of mathematics studying **decision making**
- ▶ A **game** is a system where one or more **players** choose **actions**; the combination of these choices lead to each agent receiving a **payoff**

Game theory

- ▶ A branch of mathematics studying **decision making**
- ▶ A **game** is a system where one or more **players** choose **actions**; the combination of these choices lead to each agent receiving a **payoff**
- ▶ Important applications in economics, ecology and social sciences as well as AI

The Prisoner's Student's Dilemma

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled
- ▶ If **Bob betrays Alice**, he receives an A whilst she gets expelled

The Prisoner's Student's Dilemma

- ▶ Two students, **Alice** and **Bob**, are suspected of copying from each other
- ▶ Each is offered a deal in exchange for information
- ▶ Each can choose to **betray** the other or stay **silent** — but they **cannot communicate** before deciding what to do
- ▶ If **both stay silent**, both receive a C grade
- ▶ If **Alice betrays Bob**, she receives an A whilst he gets expelled
- ▶ If **Bob betrays Alice**, he receives an A whilst she gets expelled
- ▶ If **both betray each other**, both get an F

Payoff matrix

Payoff matrix

	A silent	A betray
B silent	A: 50 B: 50	A: 70 B: -100
B betray	A: -100 B: 70	A: 0 B: 0

Alice's thought process

Alice's thought process

- The best outcome overall is for both of us to stay silent

Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
- ▶ However if Bob stays silent, I can get a better mark by betraying him

Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
- ▶ However if Bob stays silent, I can get a better mark by betraying him
- ▶ And if Bob betrays me, I should betray him to avoid getting expelled

Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
- ▶ However if Bob stays silent, I can get a better mark by betraying him
- ▶ And if Bob betrays me, I should betray him to avoid getting expelled
- ▶ Therefore the rational choice is to betray

Alice's thought process

- ▶ The best outcome overall is for both of us to stay silent
 - ▶ However if Bob stays silent, I can get a better mark by betraying him
 - ▶ And if Bob betrays me, I should betray him to avoid getting expelled
 - ▶ Therefore the rational choice is to betray
- ... and Bob's thought process is the same!

Nash equilibrium

Nash equilibrium

- ▶ Consider the situation where both have chosen to betray

Nash equilibrium

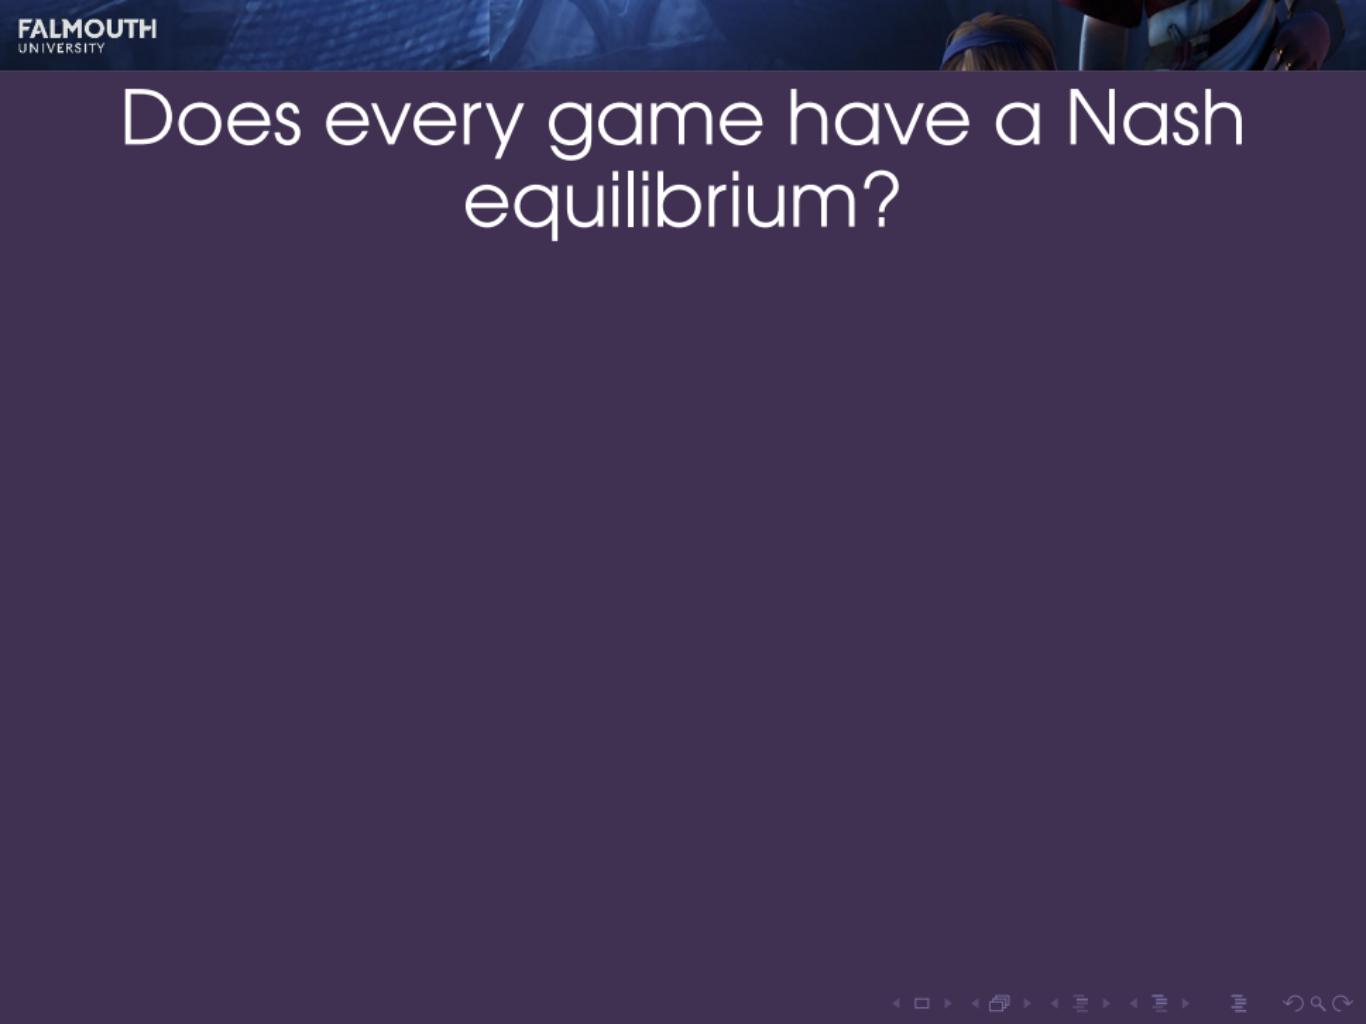
- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch

Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch
- ▶ Such a situation is called a **Nash equilibrium**

Nash equilibrium

- ▶ Consider the situation where both have chosen to betray
- ▶ Neither person has anything to gain by switching to silence, assuming the other person doesn't also switch
- ▶ Such a situation is called a **Nash equilibrium**
- ▶ If all players are **rational** (in the sense of wanting to maximising payoff), they should converge upon a Nash equilibrium



Does every game have a Nash equilibrium?

Does every game have a Nash equilibrium?

	A rock	A paper	A scissors
B rock	A: 0 B: 0	A: +1 B: -1	A: -1 B: +1
B paper	A: -1 B: +1	A: 0 B: 0	A: +1 B: -1
B scissors	A: +1 B: -1	A: -1 B: +1	A: 0 B: 0

Nash equilibrium for Rock-Paper-Scissors

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
 - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
 - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...
- ▶ The optimum strategy is to be **unpredictable**

Nash equilibrium for Rock-Paper-Scissors

- ▶ Committing to any choice of action can be **exploited**
- ▶ E.g. if you always choose paper, I choose scissors
- ▶ If we try to reason naïvely, we get stuck in a loop
 - ▶ If I choose paper, you'll choose scissors, so I should choose rock, but then you'll choose paper, so I'll choose scissors, so you'll choose rock, so I choose paper...
- ▶ The optimum strategy is to be **unpredictable**
- ▶ Choose rock with probability $\frac{1}{3}$, paper with probability $\frac{1}{3}$, scissors with probability $\frac{1}{3}$

Mixed strategies

Mixed strategies

- A **mixed strategy** assigns probabilities to actions and chooses one at random

Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random
- ▶ In contrast to a **pure** or **deterministic strategy**, which always chooses the same action

Mixed strategies

- ▶ A **mixed strategy** assigns probabilities to actions and chooses one at random
- ▶ In contrast to a **pure** or **deterministic strategy**, which always chooses the same action
- ▶ If we allow mixed strategies, **every game has at least one Nash equilibrium**



Guess $\frac{2}{3}$ of the average

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to $\frac{2}{3}$ of the mean of all guesses

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to $\frac{2}{3}$ of the mean of all guesses
- ▶ Example:

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to $\frac{2}{3}$ of the mean of all guesses
- ▶ Example:
 - ▶ If the guesses are 30, 40 and 80...

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to $\frac{2}{3}$ of the mean of all guesses
- ▶ Example:
 - ▶ If the guesses are 30, 40 and 80...
 - ▶ ... then the mean is $\frac{30+40+80}{3} = 50...$

Guess $\frac{2}{3}$ of the average

- ▶ Everyone guesses a real number (decimals are allowed) between 0 and 100 inclusive
- ▶ The winner is the person who guesses closest to $\frac{2}{3}$ of the mean of all guesses
- ▶ Example:
 - ▶ If the guesses are 30, 40 and 80...
 - ▶ ... then the mean is $\frac{30+40+80}{3} = 50...$
 - ▶ ... so the winning guess is 30, as this is closest to $\frac{2}{3} \times 50 = 33.333$

The rational guess

The rational guess

- The average can't possibly be greater than 100

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629
- ▶ ... and so on ad infinitum

The rational guess

- ▶ The average can't possibly be greater than 100
- ▶ So no rational player would guess a number greater than 66.666
- ▶ Which means the average can't possibly be greater than 66.666
- ▶ So no rational player would guess greater than 44.444
- ▶ Which means the average can't possibly be greater than 44.444
- ▶ So no rational player would guess greater than 29.629
- ▶ ... and so on ad infinitum
- ▶ So the only **rational** guess is 0, as every rational player should guess 0 and $\frac{2}{3}$ of 0 is 0

Rationality

Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers

Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers
- ▶ However it's important to remember that **humans aren't always rational**

Rationality

- ▶ Rationality is a useful assumption for mathematics and AI programmers
- ▶ However it's important to remember that **humans aren't always rational**
- ▶ E.g. in guess $\frac{2}{3}$, if enough players deviate from the rational guess of 0, then 0 is no longer the best guess!

Planning



Planning

Planning

- ▶ An **agent** in an **environment**

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**
- ▶ Problem: find a sequence of actions that leads to the goal

STRIPS planning

STRIPS planning

- Stanford Research Institute Problem Solver

STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true

STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:

STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)

STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)

STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:

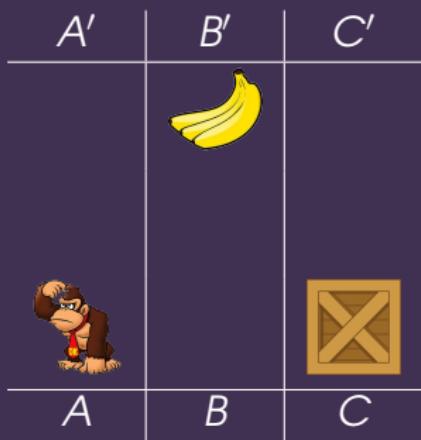
STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:
 - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)

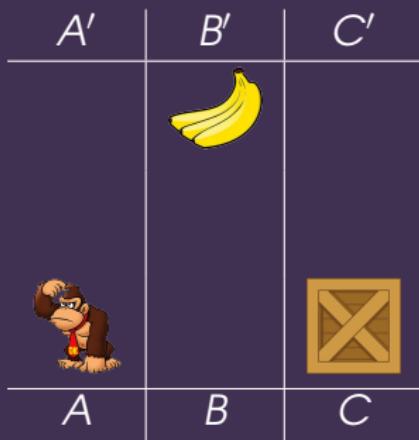
STRIPS planning

- ▶ Stanford Research Institute Problem Solver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:
 - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)
 - ▶ Postconditions (specifying what predicates are made true or false by this action)

STRIPS example



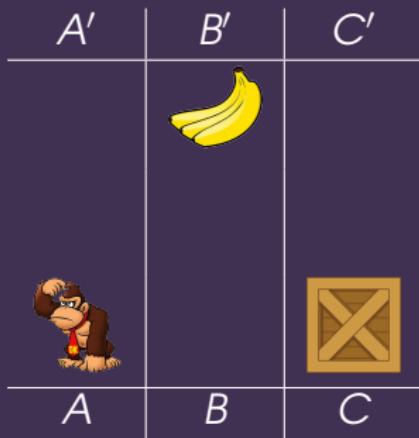
STRIPS example



Initial state:

At (A) ,
BoxAt (C) ,
BananasAt (B')

STRIPS example



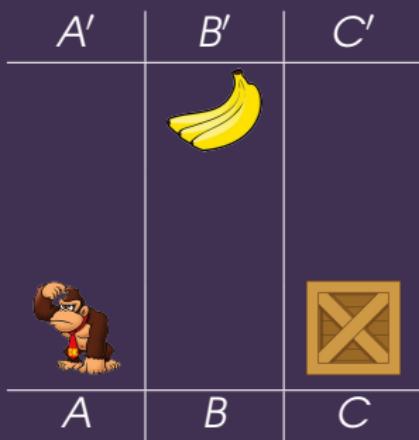
Initial state:

At (A) ,
BoxAt (C) ,
BananasAt (B')

Goal:

HasBananas

STRIPS example — Actions



Move(x , y)

Pre: At(x)

Post: !At(x), At(y)

ClimbUp(x)

Pre: At(x), BoxAt(x)

Post: !At(x), At(x')

ClimbDown(x')

Pre: At(x'), BoxAt(x)

Post: !At(x'), At(x)

PushBox(x , y)

Pre: At(x), BoxAt(x)

Post: !At(x), At(y),

!BoxAt(x), BoxAt(y)

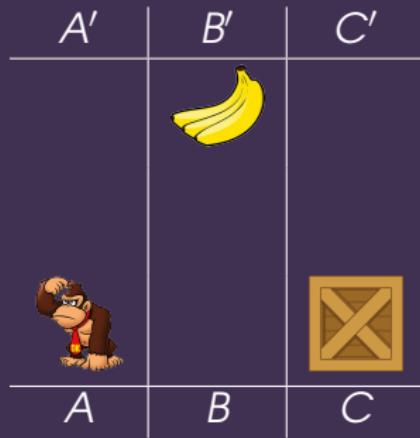
TakeBananas(x)

Pre: At(x), BananasAt(x)

Post: !BananasAt(x), HasBananas



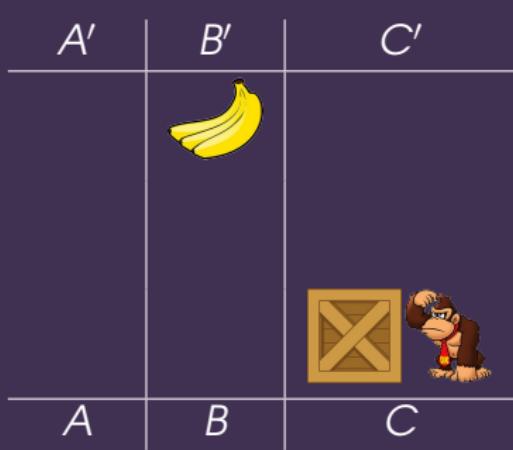
STRIPS example — Solution



State:

At (A),
BoxAt (C),
BananasAt (B')

STRIPS example — Solution



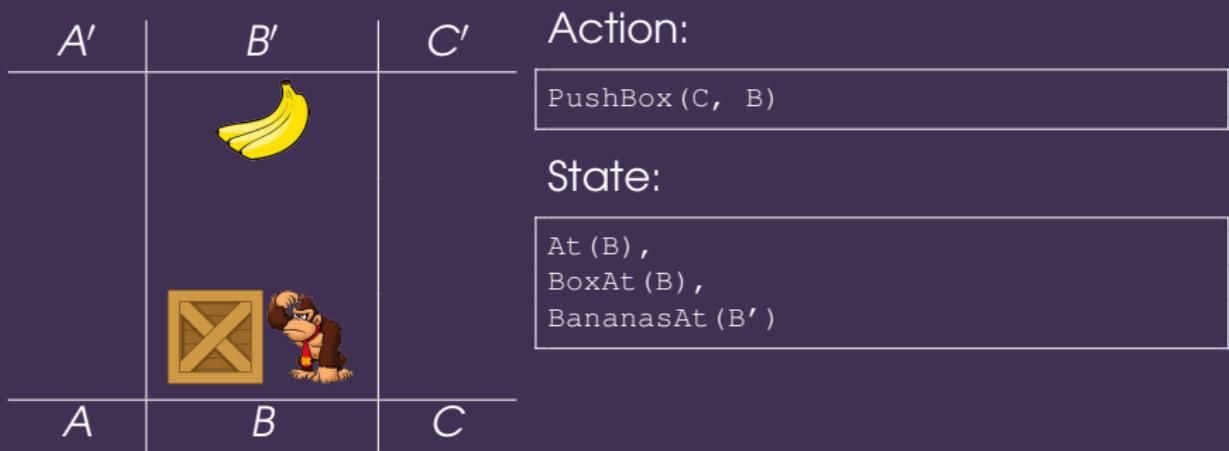
Action:

Move (A, C)

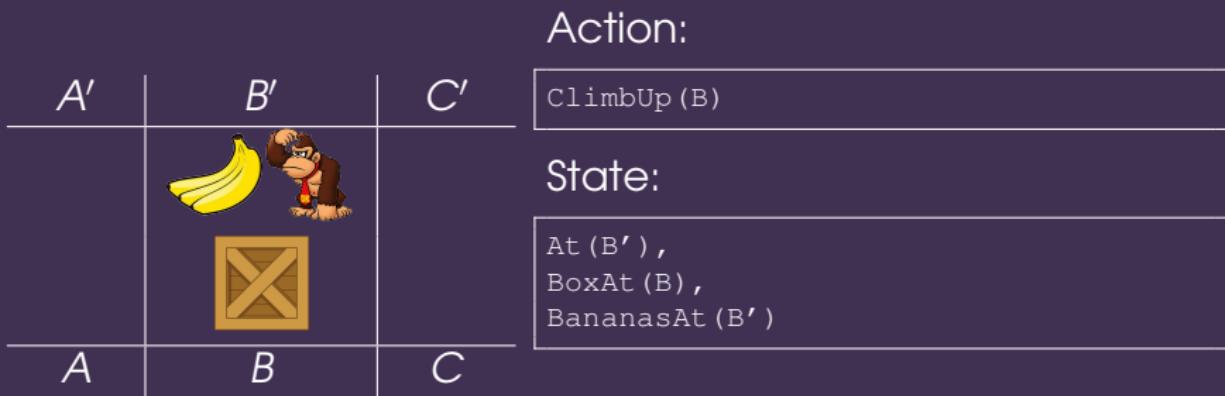
State:

At (C),
BoxAt (C),
BananasAt (B')

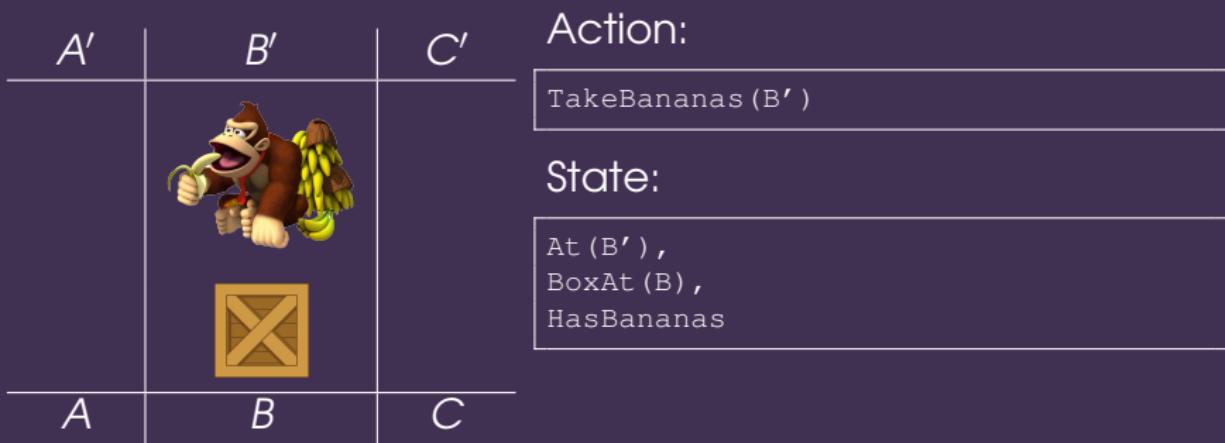
STRIPS example — Solution



STRIPS example — Solution



STRIPS example — Solution



Finding the solution

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
 - ▶ Nodes: environment states

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
 - ▶ Nodes: environment states
 - ▶ Edges: actions

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ We can construct a **state-action graph**
 - ▶ Nodes: environment states
 - ▶ Edges: actions
- ▶ We can then **search** this graph to find the optimum path from the current state to a goal state

Searching for the solution

Searching for the solution

- We have a **tree**, which is a type of **graph**

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
 - ▶ Depth-first search

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
 - ▶ Depth-first search
 - ▶ Breadth-first search

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
 - ▶ Depth-first search
 - ▶ Breadth-first search
 - ▶ Dijkstra's algorithm (same as BFS unless actions have non-uniform **cost**)

Searching for the solution

- ▶ We have a **tree**, which is a type of **graph**
- ▶ We have an **initial node** within this tree
- ▶ Want to find a **sequence of edges** that leads to a **goal node**
- ▶ Does this sound familiar?
- ▶ Very similar to **pathfinding**, so can use the same algorithms (recall from COMP280 session 8)
 - ▶ Depth-first search
 - ▶ Breadth-first search
 - ▶ Dijkstra's algorithm (same as BFS unless actions have non-uniform **cost**)
 - ▶ A* (if we have a suitable **heuristic**)

GOAP



GOAP

GOAP

- ▶ Goal Oriented Action Planning

GOAP

- ▶ **Goal Oriented Action Planning**
- ▶ Originally developed for F.E.A.R. (2005), since used in several games

GOAP

- ▶ **Goal Oriented Action Planning**
- ▶ Originally developed for F.E.A.R. (2005), since used in several games
- ▶ A modified version of STRIPS specifically for real-time planning in video games

GOAP

GOAP

- ▶ Each agent has a **goal set**

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**
 - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**
 - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**
 - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
 - ▶ Like in STRIPS — actions have preconditions and postconditions

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**
 - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
 - ▶ Like in STRIPS — actions have preconditions and postconditions
 - ▶ Unlike STRIPS, each action also has a **cost**

Action sets

Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**

Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**

Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system

Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system
- ▶ E.g. this was used by the F.E.A.R. team to quickly add new enemy types

Action sets

Soldier

- Action
 - 1 AI/Actions/Attack
 - 2 AI/Actions/AttackCrouch
 - 3 AI/Actions/SuppressionFire
 - 4 AI/Actions/SuppressionFireFromCover
 - 5 AI/Actions/FlushOutWithGrenade
 - 6 AI/Actions/AttackFromCover
 - 7 AI/Actions/BlindFireFromCover
 - 8 AI/Actions/AttackGrenadeFromCover
 - 9 AI/Actions/AttackFromView
 - 10 AI/Actions/DrawWeapon
 - 11 AI/Actions/HolsterWeapon
 - 12 AI/Actions/ReloadCrouch
 - 13 AI/Actions/ReloadCovered
 - 14 AI/Actions/InspectDisturbance
 - 15 AI/Actions/LookAtDisturbance
 - 16 AI/Actions/SurveyArea
 - 17 AI/Actions/DodgeRoll
 - 18 AI/Actions/DodgeShuffle
 - 19 AI/Actions/DodgeCovered
 - 20 AI/Actions/Uncover
 - 21 AI/Actions/AttackMelee

Assassin

- Action
 - 1 AI/Actions/Attack
 - 2 AI/Actions/InspectDisturbance
 - 3 AI/Actions/LookAtDisturbance
 - 4 AI/Actions/SurveyArea
 - 5 AI/Actions/AttackMeleeUncloaked
 - 6 AI/Actions/TraverseBlockedDoor
 - 7 AI/Actions/UseSmartObjectNodeMounted
 - 8 AI/Actions/MountNodeUncloaked
 - 9 AI/Actions/DismountNodeUncloaked
 - 10 AI/Actions/TraverseLinkUncloaked
 - 11 AI/Actions/AttackFromAmbush
 - 12 AI/Actions/DodgeRollParanoid
 - 13 AI/Actions/AttackLungeUncloaked
 - 14 AI/Actions/LopeToTargetUncloaked

Rat

- Action
 - 1 AI/Actions/Animate
 - 2 AI/Actions/Idle
 - 3 AI/Actions/GotoNode
 - 4 AI/Actions/UseSmartObjectNode

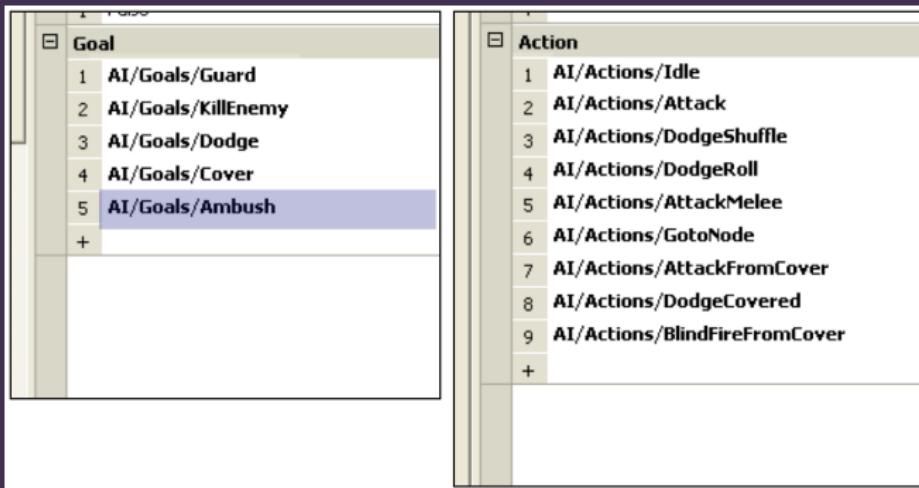
Layering

Layering

- Goal set allows different behaviours with different priorities to be **layered**

Layering

- ▶ Goal set allows different behaviours with different priorities to be **layered**
- ▶ E.g. enemy AI in F.E.A.R.:



Implementing GOAP

Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning

Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values

Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation

Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation
- ▶ Most implementations also allow for **programmatic** preconditions (e.g. calling the pathfinding system to check availability of a path)

Implementing GOAP

Implementing GOAP

- ▶ Not difficult to implement

Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist

Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist
- ▶ Not built into Unity or Unreal, but asset store packages are available

Finding the plan

Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions

Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use Dijkstra or A* to find the lowest cost path to the goal

Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use Dijkstra or A* to find the lowest cost path to the goal
- ▶ Plan is a **queue** of actions that the agent then executes

Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use Dijkstra or A* to find the lowest cost path to the goal
- ▶ Plan is a **queue** of actions that the agent then executes
- ▶ If the plan is interrupted or fails then the agent can **replan**

GOAP vs behaviour trees

GOAP vs behaviour trees

- BT: Designer specifies “how”

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**
- ▶ GOAP can be classified as **computational intelligence**