



COMP250: Artificial Intelligence

4: Utility-Based AI

Utility

Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?
- ▶ **Utility theory** supposes that a given outcome can be assigned a **single number** measuring its **utility**
- ▶ Actions can then be **ranked** by utility, and one with the **highest** utility chosen
- ▶ Utility is sometimes called **reward**, **payoff**, **fitness**
- ▶ Multiply by -1 and we have **cost**

Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Assuming everything else equal, where do you buy it from?

Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Amazon offers next-day delivery, but delivery from Bob's Computers takes 4 weeks
- ▶ Assuming everything else equal, where do you buy it from?
- ▶ Utility is a **single number** — we essentially have to put a monetary value on the longer wait time

Weighting

- ▶ Utility is often formed of several **decision factors**
- ▶ In the laptop buying example: price and delivery time
- ▶ To get a utility value we can take a **weighted sum**


```
utility = WEIGHT_PRICE * price + WEIGHT_TIME * time;
```

- ▶ Here `WEIGHT_PRICE` and `WEIGHT_TIME` are constant values
- ▶ The values used will influence the agent's behaviour and so must be carefully tuned by the designer

Nonlinearity

- ▶ Decision factors are not always linear
- ▶ E.g. a difference of £1 is more significant for something that costs £5 than something that costs £500
- ▶ E.g. a difference of 1 HP is more significant if the agent is close to death than if it is at full health
- ▶ Therefore we may want to apply a **curve** mapping to decision factors

Polynomial curve




► Formula: $y = x^k$

► C#:

`Math.Pow(x, k)`


► k is a constant:
bigger k gives a
steeper curve

Inverse polynomial curve



- ▶ Same formula as polynomial curve, but k is between 0 and 1
- ▶ k closer to 0 gives a steeper curve

Step function



► Formula:


$$y = \begin{cases} 0 & \text{if } x < 0.5 \\ 1 & \text{if } x \geq 0.5 \end{cases}$$

► C#:

`(x < 0.5f) ? 0 : 1`

► Models a **threshold**
or **if-then** rule

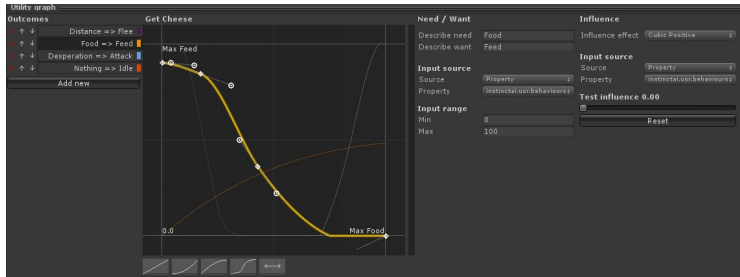
Logistic function



- ▶ Formula: $y = \frac{1}{1+e^{-kx}}$
- ▶ C#:
`1 / (1 + Mathf.Exp(-k*x))`
- ▶ Similar to step function but with a softer transition from “off” to “on”
- ▶ k is a constant: bigger k gives a steeper curve

Tweaking curves

- ▶ Adjusting utility curves is **more art than science**
- ▶ It's worth **experimenting** with different curve types to get the desired result
- ▶ **Graphical curve editors** are worth investigating
- ▶ E.g. InstinctAI asset for Unity:



Expected utility

- ▶ If the environment is **stochastic**, the outcome of an action (and hence its utility) may not be known with certainty
- ▶ Let $p(x)$ be the probability that a given action has utility x
- ▶ Then the **expected utility** is

$$\sum_x x \cdot p(x)$$

- ▶ That is, the sum of utility values weighted by their probabilities

Expected utility — example

- ▶ A slot machine pays out:
 - ▶ £1 with probability 0.05
 - ▶ £5 with probability 0.03
 - ▶ £10 with probability 0.02
 - ▶ Nothing with probability 0.9
- ▶ The expected payout is

$$1 \times 0.05 + 5 \times 0.03 + 10 \times 0.02 + 0 \times 0.9 = 0.4$$

i.e. £0.40

- ▶ If it costs £1 to play the slot machine, the expected utility overall is −£0.60
- ▶ (Although the actual utility can range from −£1 to +£9)

Inertia


- ▶ Utilities will generally change rapidly as the state of the environment changes
- ▶ This could result in the agent **changing its mind** as to which action has the best utility
- ▶ Can force the agent to finish its current action before evaluating and choosing another
- ▶ Can give a utility bonus to sticking to the current action — this still allows the agent to change its mind if the current action's utility becomes very bad

Multi-objective optimisation

- ▶ Utility-based AI is **single-objective**: all decision factors must be combined into a single number
- ▶ An alternative is **multi-objective**: treat all decision factors as separate, and find an action that optimises all of them at once

Pareto optimality

- ▶ Consider the space of all possible solutions (e.g. actions or plans)
- ▶ A solution is **Pareto dominated** if there is some other solution that is better than it on **all** decision factors at once
- ▶ A solution is **Pareto optimal** if it is not dominated by any other solution



Single-objective vs multi-objective

- ▶ Multi-objective optimisation gets around the problem of having to tune weights for decision factors
- ▶ However, there are generally a large number of Pareto optimal solutions so we need some other method to tie-break between them

GOAP

GOAP

- ▶ **Goal Oriented Action Planning**
- ▶ Originally developed for F.E.A.R. (2005), since used in several games
- ▶ A modified version of STRIPS planning (recall from last week) specifically for real-time planning in video games

GOAP

- ▶ Each agent has a **goal set**
 - ▶ Multiple goals with differing **priority**
 - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
 - ▶ Like in STRIPS — actions have preconditions and postconditions
 - ▶ Unlike STRIPS, each action also has a **cost**

Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system
- ▶ E.g. this was used by the F.E.A.R. team to quickly add new enemy types

Action sets

Soldier

Action	
1	AI/Actions/Attack
2	AI/Actions/AttackCrouch
3	AI/Actions/SuppressionFire
4	AI/Actions/SuppressionFireFromCover
5	AI/Actions/FlushOutWithGrenade
6	AI/Actions/AttackFromCover
7	AI/Actions/BlindFireFromCover
8	AI/Actions/AttackGrenadeFromCover
9	AI/Actions/AttackFromView
10	AI/Actions/DrawWeapon
11	AI/Actions/HolsterWeapon
12	AI/Actions/ReloadCrouch
13	AI/Actions/ReloadCovered
14	AI/Actions/InspectDisturbance
15	AI/Actions/LookAtDisturbance
16	AI/Actions/SurveyArea
17	AI/Actions/DodgeRoll
18	AI/Actions/DodgeShuffle
19	AI/Actions/DodgeCovered
20	AI/Actions/Uncover
21	AI/Actions/AttackMelee

Assassin

Action	
1	AI/Actions/Attack
2	AI/Actions/InspectDisturbance
3	AI/Actions/LookAtDisturbance
4	AI/Actions/SurveyArea
5	AI/Actions/AttackMeleeUncloaked
6	AI/Actions/TraverseBlockedDoor
7	AI/Actions/UseSmartObjectNodeMounted
8	AI/Actions/MountNodeUncloaked
9	AI/Actions/DismountNodeUncloaked
10	AI/Actions/TraverseLinkUncloaked
11	AI/Actions/AttackFromAmbush
12	AI/Actions/DodgeRollParanoid
13	AI/Actions/AttackLungeUncloaked
14	AI/Actions/LopeToTargetUncloaked
+	

Rat

Action	
1	AI/Actions/Animate
2	AI/Actions/Idle
3	AI/Actions/GotoNode
4	AI/Actions/UseSmartObjectNode
+	

Layering

- ▶ Goal set allows different behaviours with different priorities to be **layered**
- ▶ E.g. enemy AI in F.E.A.R.:

Goal	
1	AI/Goals/Guard
2	AI/Goals/KillEnemy
3	AI/Goals/Dodge
4	AI/Goals/Cover
5	AI/Goals/Ambush
+	

Action	
1	AI/Actions/Idle
2	AI/Actions/Attack
3	AI/Actions/DodgeShuffle
4	AI/Actions/DodgeRoll
5	AI/Actions/AttackMelee
6	AI/Actions/GotoNode
7	AI/Actions/AttackFromCover
8	AI/Actions/DodgeCovered
9	AI/Actions/BlindFireFromCover
+	

Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation
- ▶ Most implementations also allow for **programmatic** preconditions (e.g. calling the pathfinding system to check availability of a path)

Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist
- ▶ Not built into Unity or Unreal, but asset store packages are available

Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use A* to find the lowest cost path to the goal
- ▶ Plan is a **queue** of actions that the agent then executes
- ▶ If the plan is interrupted or fails then the agent can **replan**

GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**
- ▶ GOAP can be classified as **computational intelligence**