



COMP250: Artificial Intelligence

4: Utility-Based AI

Utility

Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?
- ▶ **Utility theory** supposes that a given outcome can be assigned a **single number** measuring its **utility**
- ▶ Actions can then be **ranked** by utility, and one with the **highest** utility chosen
- ▶ Utility is sometimes called **reward**, **payoff**, **fitness**
- ▶ Multiply by -1 and we have **cost**

Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Assuming everything else equal, where do you buy it from?

Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Amazon offers next-day delivery, but delivery from Bob's Computers takes 4 weeks
- ▶ Assuming everything else equal, where do you buy it from?
- ▶ Utility is a **single number** — we essentially have to put a monetary value on the longer wait time

Weighting

- ▶ Utility is often formed of several **decision factors**
- ▶ In the laptop buying example: price and delivery time
- ▶ To get a utility value we can take a **weighted sum**

```
utility = WEIGHT_PRICE * price + WEIGHT_TIME * time;
```

- ▶ Here `WEIGHT_PRICE` and `WEIGHT_TIME` are constant values
- ▶ The values used will influence the agent's behaviour and so must be carefully tuned by the designer


Utility — example

- ▶ A GPU costs £400 from Amazon or £500 from Bob's Computers
- ▶ Both offer next day delivery
- ▶ You have a moral objection to supporting large corporations, and would rather support local businesses
- ▶ Assuming everything else equal, where do you buy it from?
- ▶ To apply utility theory, we need to **quantify** everything — which may mean putting a numerical value on **intangible** things

Nonlinearity

- ▶ Decision factors are not always linear
- ▶ E.g. a difference of £1 is more significant for something that costs £5 than something that costs £500
- ▶ E.g. a difference of 1 HP is more significant if the agent is close to death than if it is at full health
- ▶ Therefore we may want to apply a **curve** mapping to decision factors

Polynomial curve




► Formula: $y = x^k$

► C#:

`Math.Pow(x, k)`


► k is a constant:
bigger k gives a
steeper curve

Inverse polynomial curve



- ▶ Same formula as polynomial curve, but k is between 0 and 1
- ▶ k closer to 0 gives a steeper curve

Step function



► Formula:


$$y = \begin{cases} 0 & \text{if } x < 0.5 \\ 1 & \text{if } x \geq 0.5 \end{cases}$$

► C#:

`(x < 0.5f) ? 0 : 1`

► Models a **threshold**
or **if-then** rule


Logistic function



- ▶ Formula: $y = \frac{1}{1+e^{-kx}}$
- ▶ C#:
`1 / (1 + Mathf.Exp(-k*x))`
- ▶ Similar to step function but with a softer transition from “off” to “on”
- ▶ k is a constant: bigger k gives a steeper curve

Tweaking curves

- ▶ Adjusting utility curves is **more art than science**
- ▶ It's worth **experimenting** with different curve types to get the desired result
- ▶ **Graphical curve editors** are worth investigating
- ▶ E.g. `AnimationCurve` in Unity, "Curve Float" in UE4
- ▶ E.g. InstinctAI asset for Unity:




Inertia

- ▶ Utilities will generally change rapidly as the state of the environment changes
- ▶ This could result in the agent **changing its mind** as to which action has the best utility
- ▶ Can force the agent to finish its current action before evaluating and choosing another
- ▶ Can give a utility bonus to sticking to the current action — this still allows the agent to change its mind if the current action's utility becomes very bad

1-ply Search

1-ply tree

- ▶ Recall from last time: in discrete planning problems or games, we can build a **state-action tree**
- ▶ Consider a tree with only one level or **ply**



1-ply search

- ▶ Suppose we have a utility measure for states
- ▶ I.e. a function which, given a state, assigns it a utility score
- ▶ Then we can “search” for the action leading to the best utility score

1-ply search

```
procedure ONEPLYSEARCH(state)  
  bestAction  $\leftarrow$  null  
  bestUtility  $\leftarrow -\infty$   
  for each valid action from state do  
    nextState  $\leftarrow$  copy of state  
    apply action to nextState  
    utility  $\leftarrow$  UTILITY(nextState)  
    if utility > bestUtility then  
      bestAction  $\leftarrow$  action  
      bestUtility  $\leftarrow$  utility  
    end if  
  end for  
  return bestAction  
end procedure
```

Why 1-ply search?

- ▶ It is often easier to come up with a utility measure for states than for actions
 - ▶ E.g. distance to goal — greedy search
 - ▶ E.g. material evaluation in chess
- ▶ Much faster than a full-blown BFS or game tree search, if depth of forward planning is not required

Expectation

Expected utility

- ▶ If the environment is **stochastic**, the outcome of an action (and hence its utility) may not be known with certainty
- ▶ Let $p(x)$ be the probability that a given action has utility x
- ▶ Then the **expected utility** is

$$\sum_x x \cdot p(x)$$

- ▶ That is, the sum of utility values weighted by their probabilities

Expected utility — example

- ▶ A slot machine pays out:
 - ▶ £1 with probability 0.05
 - ▶ £5 with probability 0.03
 - ▶ £10 with probability 0.02
 - ▶ Nothing with probability 0.9
- ▶ The expected payout is

$$1 \times 0.05 + 5 \times 0.03 + 10 \times 0.02 + 0 \times 0.9 = 0.4$$

i.e. £0.40

- ▶ If it costs £1 to play the slot machine, the expected utility overall is −£0.60
- ▶ (Although the actual utility can range from −£1 to +£9)


Multi-Objective Optimisation

Multi-objective optimisation

- ▶ Utility-based AI is **single-objective**: all decision factors must be combined into a single number
- ▶ An alternative is **multi-objective**: treat all decision factors as separate, and find an action that optimises all of them at once

Pareto optimality

- ▶ Consider the space of all possible solutions (e.g. actions or plans)
- ▶ A solution is **Pareto dominated** if there is some other solution that is better than it on **all** decision factors at once
- ▶ A solution is **Pareto optimal** if it is not dominated by any other solution



Single-objective vs multi-objective

- ▶ Multi-objective optimisation gets around the problem of having to tune weights for decision factors
- ▶ However, there are generally a large number of Pareto optimal solutions so we need some other method to tie-break between them
- ▶ ... which may boil down to weights (or at least priorities) anyway