



## COMP250: Artificial Intelligence

# 4: Utility-Based AI

# Utility



# Utility theory

# Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?

# Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?
- ▶ **Utility theory** supposes that a given outcome can be assigned a **single number** measuring its **utility**

# Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?
- ▶ **Utility theory** supposes that a given outcome can be assigned a **single number** measuring its **utility**
- ▶ Actions can then be **ranked** by utility, and one with the **highest** utility chosen

# Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?
- ▶ **Utility theory** supposes that a given outcome can be assigned a **single number** measuring its **utility**
- ▶ Actions can then be **ranked** by utility, and one with the **highest** utility chosen
- ▶ Utility is sometimes called **reward**, **payoff**, **fitness**

# Utility theory

- ▶ How does an AI agent measure “goodness” or “usefulness” of the outcomes of its actions?
- ▶ **Utility theory** supposes that a given outcome can be assigned a **single number** measuring its **utility**
- ▶ Actions can then be **ranked** by utility, and one with the **highest** utility chosen
- ▶ Utility is sometimes called **reward**, **payoff**, **fitness**
- ▶ Multiply by  $-1$  and we have **cost**



# Utility — example

# Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers

# Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Assuming everything else equal, where do you buy it from?

# Utility — example

# Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers

# Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Amazon offers next-day delivery, but delivery from Bob's Computers takes 4 weeks

# Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Amazon offers next-day delivery, but delivery from Bob's Computers takes 4 weeks
- ▶ Assuming everything else equal, where do you buy it from?

# Utility — example

- ▶ A laptop costs £500 from Amazon or £450 from Bob's Computers
- ▶ Amazon offers next-day delivery, but delivery from Bob's Computers takes 4 weeks
- ▶ Assuming everything else equal, where do you buy it from?
- ▶ Utility is a **single number** — we essentially have to put a monetary value on the longer wait time



# Weighting

# Weighting

- ▶ Utility is often formed of several **decision factors**

# Weighting

- ▶ Utility is often formed of several **decision factors**
- ▶ In the laptop buying example: price and delivery time

# Weighting

- ▶ Utility is often formed of several **decision factors**
- ▶ In the laptop buying example: price and delivery time
- ▶ To get a utility value we can take a **weighted sum**

```
utility = WEIGHT_PRICE * price + WEIGHT_TIME * time;
```

# Weighting

- ▶ Utility is often formed of several **decision factors**
- ▶ In the laptop buying example: price and delivery time
- ▶ To get a utility value we can take a **weighted sum**

```
utility = WEIGHT_PRICE * price + WEIGHT_TIME * time;
```

- ▶ Here `WEIGHT_PRICE` and `WEIGHT_TIME` are constant values

# Weighting

- ▶ Utility is often formed of several **decision factors**
- ▶ In the laptop buying example: price and delivery time
- ▶ To get a utility value we can take a **weighted sum**

```
utility = WEIGHT_PRICE * price + WEIGHT_TIME * time;
```

- ▶ Here `WEIGHT_PRICE` and `WEIGHT_TIME` are constant values
- ▶ The values used will influence the agent's behaviour and so must be carefully tuned by the designer

# Nonlinearity

# Nonlinearity

- ▶ Decision factors are not always linear



# Nonlinearity

- ▶ Decision factors are not always linear
- ▶ E.g. a difference of £1 is more significant for something that costs £5 than something that costs £500


# Nonlinearity

- ▶ Decision factors are not always linear
- ▶ E.g. a difference of £1 is more significant for something that costs £5 than something that costs £500
- ▶ E.g. a difference of 1 HP is more significant if the agent is close to death than if it is at full health


# Nonlinearity

- ▶ Decision factors are not always linear
- ▶ E.g. a difference of £1 is more significant for something that costs £5 than something that costs £500
- ▶ E.g. a difference of 1 HP is more significant if the agent is close to death than if it is at full health
- ▶ Therefore we may want to apply a **curve** mapping to decision factors

# Polynomial curve




# Polynomial curve



► Formula:  $y = x^k$

# Polynomial curve




► Formula:  $y = x^k$

► C#:


`Math.Pow(x, k)`

# Polynomial curve




- ▶ Formula:  $y = x^k$
- ▶ C#:  
`Math.Pow(x, k)`
- ▶  $k$  is a constant:  
bigger  $k$  gives a  
steeper curve

# Inverse polynomial curve






# Inverse polynomial curve




- Same formula as polynomial curve, but  $k$  is between 0 and 1

# Inverse polynomial curve




- ▶ Same formula as polynomial curve, but  $k$  is between 0 and 1
- ▶  $k$  closer to 0 gives a steeper curve

# Step function




# Step function



► Formula:

$$y = \begin{cases} 0 & \text{if } x < 0.5 \\ 1 & \text{if } x \geq 0.5 \end{cases}$$

# Step function



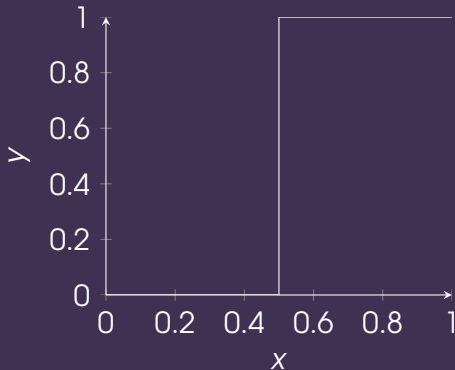
► Formula:

$$y = \begin{cases} 0 & \text{if } x < 0.5 \\ 1 & \text{if } x \geq 0.5 \end{cases}$$

► C#:

`(x < 0.5f) ? 0 : 1`

# Step function



- Formula:


$$y = \begin{cases} 0 & \text{if } x < 0.5 \\ 1 & \text{if } x \geq 0.5 \end{cases}$$

- C#:


`(x < 0.5f) ? 0 : 1`

- Models a **threshold** or **if-then** rule

# Logistic function




# Logistic function



► Formula:  $y = \frac{1}{1+e^{-kx}}$



# Logistic function




► Formula:  $y = \frac{1}{1+e^{-kx}}$

► C#:


```
1 / (1 + Mathf.Exp(-k*x))
```

# Logistic function



- ▶ Formula:  $y = \frac{1}{1+e^{-kx}}$
- ▶ C#:  
`1 / (1 + Mathf.Exp(-k*x))`
- ▶ Similar to step function but with a softer transition from “off” to “on”

# Logistic function



- ▶ Formula:  $y = \frac{1}{1+e^{-kx}}$
- ▶ C#:  
`1 / (1 + Mathf.Exp(-k*x))`
- ▶ Similar to step function but with a softer transition from “off” to “on”
- ▶  $k$  is a constant: bigger  $k$  gives a steeper curve

# Tweaking curves

# Tweaking curves

- ▶ Adjusting utility curves is **more art than science**

# Tweaking curves

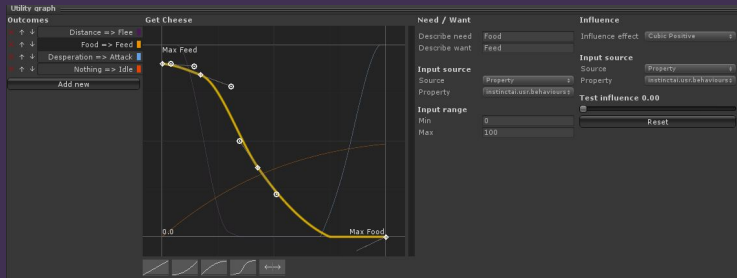
- ▶ Adjusting utility curves is **more art than science**
- ▶ It's worth **experimenting** with different curve types to get the desired result

# Tweaking curves

- ▶ Adjusting utility curves is **more art than science**
- ▶ It's worth **experimenting** with different curve types to get the desired result
- ▶ **Graphical curve editors** are worth investigating

# Tweaking curves

- ▶ Adjusting utility curves is **more art than science**
- ▶ It's worth **experimenting** with different curve types to get the desired result
- ▶ **Graphical curve editors** are worth investigating
- ▶ E.g. InstinctAI asset for Unity:





# Expected utility

# Expected utility

- ▶ If the environment is **stochastic**, the outcome of an action (and hence its utility) may not be known with certainty

# Expected utility

- ▶ If the environment is **stochastic**, the outcome of an action (and hence its utility) may not be known with certainty
- ▶ Let  $p(x)$  be the probability that a given action has utility  $x$

# Expected utility

- ▶ If the environment is **stochastic**, the outcome of an action (and hence its utility) may not be known with certainty
- ▶ Let  $p(x)$  be the probability that a given action has utility  $x$
- ▶ Then the **expected utility** is

$$\sum_x x \cdot p(x)$$

# Expected utility

- ▶ If the environment is **stochastic**, the outcome of an action (and hence its utility) may not be known with certainty
- ▶ Let  $p(x)$  be the probability that a given action has utility  $x$
- ▶ Then the **expected utility** is

$$\sum_x x \cdot p(x)$$

- ▶ That is, the sum of utility values weighted by their probabilities

# Expected utility — example

# Expected utility — example

- ▶ A slot machine pays out:

# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05



# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05
  - ▶ £5 with probability 0.03

# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05
  - ▶ £5 with probability 0.03
  - ▶ £10 with probability 0.02

# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05
  - ▶ £5 with probability 0.03
  - ▶ £10 with probability 0.02
  - ▶ Nothing with probability 0.9

# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05
  - ▶ £5 with probability 0.03
  - ▶ £10 with probability 0.02
  - ▶ Nothing with probability 0.9
- ▶ The expected payout is

$$1 \times 0.05 + 5 \times 0.03 + 10 \times 0.02 + 0 \times 0.9 = 0.4$$

i.e. £0.40

# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05
  - ▶ £5 with probability 0.03
  - ▶ £10 with probability 0.02
  - ▶ Nothing with probability 0.9
- ▶ The expected payout is

$$1 \times 0.05 + 5 \times 0.03 + 10 \times 0.02 + 0 \times 0.9 = 0.4$$

i.e. £0.40

- ▶ If it costs £1 to play the slot machine, the expected utility overall is –£0.60

# Expected utility — example

- ▶ A slot machine pays out:
  - ▶ £1 with probability 0.05
  - ▶ £5 with probability 0.03
  - ▶ £10 with probability 0.02
  - ▶ Nothing with probability 0.9
- ▶ The expected payout is

$$1 \times 0.05 + 5 \times 0.03 + 10 \times 0.02 + 0 \times 0.9 = 0.4$$

i.e. £0.40

- ▶ If it costs £1 to play the slot machine, the expected utility overall is –£0.60
- ▶ (Although the actual utility can range from –£1 to +£9)

# Inertia

# Inertia

- ▶ Utilities will generally change rapidly as the state of the environment changes



# Inertia

- ▶ Utilities will generally change rapidly as the state of the environment changes
- ▶ This could result in the agent **changing its mind** as to which action has the best utility

# Inertia

- ▶ Utilities will generally change rapidly as the state of the environment changes
- ▶ This could result in the agent **changing its mind** as to which action has the best utility
- ▶ Can force the agent to finish its current action before evaluating and choosing another

# Inertia

- ▶ Utilities will generally change rapidly as the state of the environment changes
- ▶ This could result in the agent **changing its mind** as to which action has the best utility
- ▶ Can force the agent to finish its current action before evaluating and choosing another
- ▶ Can give a utility bonus to sticking to the current action — this still allows the agent to change its mind if the current action's utility becomes very bad

# Multi-objective optimisation

# Multi-objective optimisation

- ▶ Utility-based AI is **single-objective**: all decision factors must be combined into a single number

# Multi-objective optimisation

- ▶ Utility-based AI is **single-objective**: all decision factors must be combined into a single number
- ▶ An alternative is **multi-objective**: treat all decision factors as separate, and find an action that optimises all of them at once

# Pareto optimality

# Pareto optimality

- ▶ Consider the space of all possible solutions (e.g. actions or plans)




# Pareto optimality

- ▶ Consider the space of all possible solutions (e.g. actions or plans)
- ▶ A solution is **Pareto dominated** if there is some other solution that is better than it on **all** decision factors at once

# Pareto optimality

- ▶ Consider the space of all possible solutions (e.g. actions or plans)
- ▶ A solution is **Pareto dominated** if there is some other solution that is better than it on **all** decision factors at once
- ▶ A solution is **Pareto optimal** if it is not dominated by any other solution



# Single-objective vs multi-objective

# Single-objective vs multi-objective

- ▶ Multi-objective optimisation gets around the problem of having to tune weights for decision factors

# Single-objective vs multi-objective

- ▶ Multi-objective optimisation gets around the problem of having to tune weights for decision factors
- ▶ However, there are generally a large number of Pareto optimal solutions so we need some other method to tie-break between them

GOAP



# GOAP

# GOAP

## ► Goal Oriented Action Planning



# GOAP

- ▶ **G**oal **O**riented **A**ction **P**lanning
- ▶ Originally developed for F.E.A.R. (2005), since used in several games

# GOAP

- ▶ **G**oal **O**riented **A**ction **P**lanning
- ▶ Originally developed for F.E.A.R. (2005), since used in several games
- ▶ A modified version of STRIPS planning (recall from last week) specifically for real-time planning in video games

# GOAP

# GOAP

- ▶ Each agent has a **goal set**

# GOAP

- ▶ Each agent has a **goal set**
  - ▶ Multiple goals with differing **priority**

# GOAP

- ▶ Each agent has a **goal set**
  - ▶ Multiple goals with differing **priority**
  - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy

# GOAP

- ▶ Each agent has a **goal set**
  - ▶ Multiple goals with differing **priority**
  - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**

# GOAP

- ▶ Each agent has a **goal set**
  - ▶ Multiple goals with differing **priority**
  - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
  - ▶ Like in STRIPS — actions have preconditions and postconditions



# GOAP

- ▶ Each agent has a **goal set**
  - ▶ Multiple goals with differing **priority**
  - ▶ Goals are like in STRIPS — sets of predicates that the agent wants to satisfy
- ▶ Each agent also has a set of **actions**
  - ▶ Like in STRIPS — actions have preconditions and postconditions
  - ▶ Unlike STRIPS, each action also has a **cost**

# Action sets

# Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**

# Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**

# Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system

# Action sets

- ▶ Different types of agent could have the **same goals** but **different action sets**
- ▶ This will result in those agents achieving those goals in **different ways**
- ▶ NB this doesn't have to be explicitly coded — it **emerges** from the GOAP system
- ▶ E.g. this was used by the F.E.A.R. team to quickly add new enemy types

# Action sets

## Soldier

+
Action
1 AI/Actions/Attack
2 AI/Actions/AttackCrouch
3 AI/Actions/SuppressionFire
4 AI/Actions/SuppressionFireFromCover
5 AI/Actions/FlushOutWithGrenade
6 AI/Actions/AttackFromCover
7 AI/Actions/BlindFireFromCover
8 AI/Actions/AttackGrenadeFromCover
9 AI/Actions/AttackFromView
10 AI/Actions/DrawWeapon
11 AI/Actions/HolsterWeapon
12 AI/Actions/ReloadCrouch
13 AI/Actions/ReloadCovered
14 AI/Actions/InspectDisturbance
15 AI/Actions/LookAtDisturbance
16 AI/Actions/SurveyArea
17 AI/Actions/DodgeRoll
18 AI/Actions/DodgeShuffle
19 AI/Actions/DodgeCovered
20 AI/Actions/Uncover
21 AI/Actions/AttackMelee

## Assassin

+
Action
1 AI/Actions/Attack
2 AI/Actions/InspectDisturbance
3 AI/Actions/LookAtDisturbance
4 AI/Actions/SurveyArea
5 AI/Actions/AttackMeleeUncloaked
6 AI/Actions/TraverseBlockedDoor
7 AI/Actions/UseSmartObjectNodeMounted
8 AI/Actions/MountNodeUncloaked
9 AI/Actions/DismountNodeUncloaked
10 AI/Actions/TraverseLinkUncloaked
11 AI/Actions/AttackFromAmbush
12 AI/Actions/DodgeRollParanoid
13 AI/Actions/AttackLungeUncloaked
14 AI/Actions/LopeToTargetUncloaked
+

## Rat

+
Action
1 AI/Actions/Animate
2 AI/Actions/Idle
3 AI/Actions/GotoNode
4 AI/Actions/UseSmartObjectNode
+

# Layering



# Layering

- ▶ Goal set allows different behaviours with different priorities to be **layered**

# Layering

- ▶ Goal set allows different behaviours with different priorities to be **layered**
- ▶ E.g. enemy AI in F.E.A.R.:

Goal	Action
1 AI/Goals/Guard	1 AI/Actions/Idle
2 AI/Goals/KillEnemy	2 AI/Actions/Attack
3 AI/Goals/Dodge	3 AI/Actions/DodgeShuffle
4 AI/Goals/Cover	4 AI/Actions/DodgeRoll
5 AI/Goals/Ambush	5 AI/Actions/AttackMelee
+	6 AI/Actions/GotoNode
	7 AI/Actions/AttackFromCover
	8 AI/Actions/DodgeCovered
	9 AI/Actions/BlindFireFromCover
	+

# Implementing GOAP

# Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning

# Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values

# Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation

# Implementing GOAP

- ▶ An **abstracted** view of the game world is used for planning
- ▶ Represented as a fixed-length array (or struct) of values
- ▶ Predicates (preconditions, postconditions, goals) represented in terms of this array representation
- ▶ Most implementations also allow for **programmatic** preconditions (e.g. calling the pathfinding system to check availability of a path)

# Implementing GOAP



# Implementing GOAP

- ▶ Not difficult to implement

# Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist

# Implementing GOAP

- ▶ Not difficult to implement
- ▶ Open-source implementations do exist
- ▶ Not built into Unity or Unreal, but asset store packages are available

# Finding the plan

# Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions

# Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use  $A^*$  to find the lowest cost path to the goal

# Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use  $A^*$  to find the lowest cost path to the goal
- ▶ Plan is a **queue** of actions that the agent then executes

# Finding the plan

- ▶ As in STRIPS, we can build a **tree** whose nodes are world states and edges are available actions
- ▶ Since actions have costs, we can use  $A^*$  to find the lowest cost path to the goal
- ▶ Plan is a **queue** of actions that the agent then executes
- ▶ If the plan is interrupted or fails then the agent can **replan**



# GOAP vs behaviour trees

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**

# GOAP vs behaviour trees

- ▶ BT: Designer specifies “how”
- ▶ GOAP: Designer specifies “what” — “how” is in whatever system is used to implement actions (FSMs in F.E.A.R.; could use BTs or hand coding)
- ▶ Both: actions (tasks in BT) are modular and reusable between agents
- ▶ GOAP: goals are also modular and reusable
- ▶ BT: goals are not represented explicitly
- ▶ BT can be classified as **authored behaviour**
- ▶ GOAP can be classified as **computational intelligence**