

COMP250: Artificial Intelligence

4: Minimax Search

Noughts and Crosses




- ▶ Clone the following repository:
`https://github.com/Falmouth-Games-Academy/comp250-live-coding`
- ▶ Open `COMP250/04_minimax` in PyCharm and run `oxo_main.py`
- ▶ If PyCharm asks for license server information, enter `http://trlicefal.fal.ac.uk`

Minimax search



Game trees



Minimax

Minimax

- ▶ Terminal game states have a **value**

Minimax

- ▶ Terminal game states have a **value**
 - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw

Minimax

- ▶ Terminal game states have a **value**
 - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value

Minimax

- ▶ Terminal game states have a **value**
 - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value

Minimax

- ▶ Terminal game states have a **value**
 - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value
- ▶ Therefore I want to **maximise** the **minimum** value my opponent can achieve

Minimax

- ▶ Terminal game states have a **value**
 - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value
- ▶ Therefore I want to **maximise** the **minimum** value my opponent can achieve
- ▶ This is generally only true for **two-player zero-sum** games

Minimax search

Minimax search

- ▶ Recursively defines a **value** for non-terminal game states

Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move

Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move
- ▶ If it's my turn, the value is the **maximum** value over next states

Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move
- ▶ If it's my turn, the value is the **maximum** value over next states
- ▶ If it's my opponent's turn, the value is the **minimum** value over next states

Minimax search pseudocode

procedure MINIMAX(state, currentPlayer)

Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
```

Minimax search pseudocode

procedure MINIMAX(state, currentPlayer)

if state is terminal **then**

return value of state

Minimax search pseudocode

procedure MINIMAX(state, currentPlayer)

if state is terminal **then**

return value of state

else if currentPlayer = 1 **then**

Minimax search pseudocode

procedure MINIMAX(state, currentPlayer)

if state is terminal **then**

return value of state

else if currentPlayer = 1 **then**

 bestValue = $-\infty$

Minimax search pseudocode

procedure MINIMAX(state, currentPlayer)

if state is terminal **then**

return value of state

else if currentPlayer = 1 **then**

 bestValue = $-\infty$

for each possible nextState **do**

Minimax search pseudocode

procedure MINIMAX(state, currentPlayer)

if state is terminal **then**

return value of state

else if currentPlayer = 1 **then**

 bestValue = $-\infty$

for each possible nextState **do**

$v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$

Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue, v)
```

Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue, v)
    end for
  return bestValue
```

Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue, v)
    end for
    return bestValue
  else if currentPlayer = 2 then
```

Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue,  $v$ )
    end for
    return bestValue
  else if currentPlayer = 2 then
    bestValue =  $+\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MIN(bestValue,  $v$ )
    end for
    return bestValue
```

Minimax search pseudocode

```
procedure MINIMAX(state, currentPlayer)
  if state is terminal then
    return value of state
  else if currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MAX(bestValue,  $v$ )
    end for
    return bestValue
  else if currentPlayer = 2 then
    bestValue =  $+\infty$ 
    for each possible nextState do
       $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$ 
      bestValue = MIN(bestValue,  $v$ )
    end for
    return bestValue
  end if
end procedure
```

Stopping early

```
for each possible nextState do  
     $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$   
     $\text{bestValue} = \text{MAX}(\text{bestValue}, v)$   
end for
```

Stopping early

```
for each possible nextState do  
     $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$   
     $\text{bestValue} = \text{MAX}(\text{bestValue}, v)$   
end for
```

- State values are always between -1 and $+1$

Stopping early

for each possible nextState **do**

$v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$

$\text{bestValue} = \text{MAX}(\text{bestValue}, v)$

end for

- ▶ State values are always between -1 and $+1$
- ▶ So if we ever have $\text{bestValue} = 1$, we can stop early

Stopping early

```
for each possible nextState do  
     $v = \text{MINIMAX}(\text{nextState}, 3 - \text{currentPlayer})$   
     $\text{bestValue} = \text{MAX}(\text{bestValue}, v)$   
end for
```

- ▶ State values are always between -1 and $+1$
- ▶ So if we ever have $\text{bestValue} = 1$, we can stop early
- ▶ Similarly when minimising if $\text{bestValue} = -1$

Using minimax search

Using minimax search

- ▶ To decide what move to play next...

Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move

Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move
- ▶ Choose the move with the maximum score

Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move
- ▶ Choose the move with the maximum score
- ▶ If there are several with the same score, choose one at random

Minimax and game theory

Minimax and game theory

- For a **two-player zero-sum** game with **perfect information** and **sequential moves**

Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**

Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**
- ▶ I.e. a minimax player plays **perfectly**

Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**
- ▶ I.e. a minimax player plays **perfectly**
- ▶ **But...**

Heuristics for search



Minimax for larger games

Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states

Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully

Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully
 - ▶ Connect 4 has $\approx 10^{13}$ states

Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully
 - ▶ Connect 4 has $\approx 10^{13}$ states
 - ▶ Chess has $\approx 10^{47}$ states

Depth limiting

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as $+1 / 0 / -1$

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as $+1 / 0 / -1$
- ▶ For nonterminal states at depth d , apply a heuristic evaluation instead of searching deeper

Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as $+1$ / 0 / -1
- ▶ For nonterminal states at depth d , apply a heuristic evaluation instead of searching deeper
- ▶ Evaluation is a number between -1 and $+1$, estimating the probable outcome of the game

1-ply search

1-ply search

- ▶ Case $d = 1$

1-ply search

- ▶ Case $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move

1-ply search

- ▶ Case $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move
- ▶ This is computationally fast

1-ply search

- ▶ Case $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move
- ▶ This is computationally fast
- ▶ Often easier to design a “which state is better” heuristic than to directly design a “which move to play” heuristic

Move ordering

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search

Move ordering

- ▶ Minimax can **stop early** if it sees a value of $+1$ for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search
- ▶ How do we know which moves are “best” and “worst”? Use a heuristic!

Designing heuristics

Designing heuristics

- ▶ The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**

Designing heuristics

- ▶ The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**
- ▶ Good heuristic design requires **in-depth knowledge** of the tactics and strategy of the game

Designing heuristics

- ▶ The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**
- ▶ Good heuristic design requires **in-depth knowledge** of the tactics and strategy of the game
- ▶ Next time we will look at what we can do if we don't possess such knowledge

Planning



Planning

Planning

- ▶ An **agent** in an **environment**

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**

Planning

- ▶ An **agent** in an **environment**
- ▶ The environment has a **state**
- ▶ The agent can perform **actions** to change the state
- ▶ The agent wants to change the state so as to achieve a **goal**
- ▶ Problem: find a sequence of actions that leads to the goal

STRIPS planning

STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver

STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true

STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:

STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)

STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)

STRIPS planning

- ▶ **S**tanford **R**esearch Institute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:


STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:
 - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)

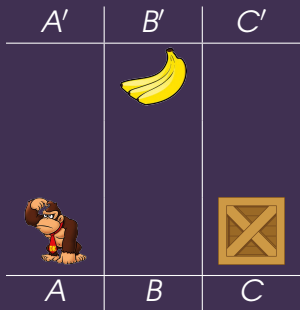
STRIPS planning

- ▶ **S**tanford **R**esearch **I**nstitute **P**roblem **S**olver
- ▶ Describes the state of the environment by a set of **predicates** which are true
- ▶ Models a problem as:
 - ▶ The **initial state** (a set of predicates which are true)
 - ▶ The **goal state** (a set of predicates, specifying whether each should be true or false)
 - ▶ The set of **actions**, each specifying:
 - ▶ Preconditions (a set of predicates which must be satisfied for this action to be possible)
 - ▶ Postconditions (specifying what predicates are made true or false by this action)

STRIPS example



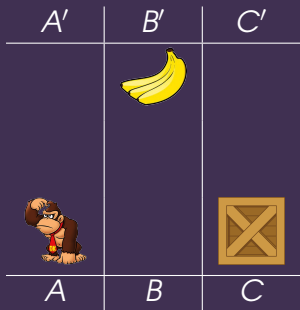
STRIPS example



Initial state:

At (A) ,
BoxAt (C) ,
BananasAt (B')

STRIPS example




Initial state:

At (A) ,
BoxAt (C) ,
BananasAt (B')

Goal:

HasBananas

STRIPS example — Actions



Move(x, y)

Pre: At(x)

Post: !At(x), At(y)

ClimbUp(x)

Pre: At(x), BoxAt(x)

Post: !At(x), At(x')

ClimbDown(x')

Pre: At(x'), BoxAt(x)

Post: !At(x'), At(x)

PushBox(x, y)

Pre: At(x), BoxAt(x)


Post: !At(x), At(y),
!BoxAt(x), BoxAt(y)

TakeBananas(x)


Pre: At(x), BananasAt(x)

Post: !BananasAt(x), HasBananas


STRIPS example — Solution




STRIPS example — Solution




STRIPS example — Solution



STRIPS example — Solution



STRIPS example — Solution



Finding the solution

Finding the solution

- For a given state, we can construct a list of all **valid actions** based on their **preconditions**

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ This should sound familiar (from 2 weeks ago)...

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ This should sound familiar (from 2 weeks ago)...
- ▶ We can construct a **tree** of states and actions

Finding the solution

- ▶ For a given state, we can construct a list of all **valid actions** based on their **preconditions**
- ▶ We can also find the **next state** resulting from each action based on their **postconditions**
- ▶ This should sound familiar (from 2 weeks ago)...
- ▶ We can construct a **tree** of states and actions
- ▶ We can then **search** this tree to find a goal state

Tree traversal

Tree traversal

procedure DEPTHFIRSTSEARCH

Tree traversal

procedure DEPTHFIRSTSEARCH
 let S be a stack

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

 let Q be a queue

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

 let Q be a queue

 enqueue root node into Q

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

 let Q be a queue

 enqueue root node into Q

while Q is not empty **do**

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

 let Q be a queue

 enqueue root node into Q

while Q is not empty **do**

 dequeue n from Q

Tree traversal

procedure DEPTHFIRSTSEARCH

 let S be a stack

 push root node onto S

while S is not empty **do**

 pop n from S

 push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

 let Q be a queue

 enqueue root node into Q

while Q is not empty **do**

 dequeue n from Q

 enqueue children of n into Q

Tree traversal

procedure DEPTHFIRSTSEARCH

let S be a stack

push root node onto S

while S is not empty **do**

pop n from S

push children of n onto S

end while

end procedure

procedure BREADTHFIRSTSEARCH

let Q be a queue

enqueue root node into Q

while Q is not empty **do**

dequeue n from Q

enqueue children of n into Q

end while

end procedure

Tree traversal example

