



COMP250: Artificial Intelligence

## **5: Game Tree Search**

# Minimax search

# Minimax

- ▶ Terminal game states have a **value**
  - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value
- ▶ Therefore I want to **maximise** the **minimum** value my opponent can achieve
- ▶ This is generally only true for **two-player zero-sum** games

# Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move
- ▶ If it's my turn, the value is the **maximum** value over next states
- ▶ If it's my opponent's turn, the value is the **minimum** value over next states

# Minimax search pseudocode

```
procedure MINIMAX(state)
  if state is terminal then
    return value of state
  else if state.currentPlayer = 1 then
    bestValue =  $-\infty$ 
    for each possible nextState do
      v = MINIMAX(nextState)
      bestValue = MAX(bestValue, v)
    return bestValue
  else if state.currentPlayer = 2 then
    bestValue =  $+\infty$ 
    for each possible nextState do
      v = MINIMAX(nextState)
      bestValue = MIN(bestValue, v)
    return bestValue
```

# Stopping early

**for each** possible nextState **do**

$v = \text{MINIMAX}(\text{nextState})$

$\text{bestValue} = \text{MAX}(\text{bestValue}, v)$

- ▶ State values are always between  $-1$  and  $+1$
- ▶ So if we ever have  $\text{bestValue} = 1$ , we can stop early
- ▶ Similarly when minimising if  $\text{bestValue} = -1$

# Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move
- ▶ Choose the move with the maximum score
- ▶ If there are several with the same score, choose one at random

# Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**
- ▶ I.e. a minimax player plays **perfectly**
- ▶ **But...**



# Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully
  - ▶ Connect 4 has  $\approx 10^{13}$  states
  - ▶ Chess has  $\approx 10^{47}$  states

**Heuristics for search**

# Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as  $+1 / 0 / -1$
- ▶ For nonterminal states at depth  $d$ , apply a heuristic evaluation instead of searching deeper
- ▶ Evaluation is a number between  $-1$  and  $+1$ , estimating the probable outcome of the game

# 1-ply search

- ▶ Case  $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move
- ▶ This is computationally fast
- ▶ Often easier to design a “which state is better” heuristic than to directly design a “which move to play” heuristic
- ▶ This is essentially a **utility-based AI**

# Move ordering

- ▶ Minimax can **stop early** if it sees a value of  $+1$  for maximising player or  $-1$  for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search
- ▶ How do we know which moves are “best” and “worst”? Use a heuristic!

# Designing heuristics

- ▶ The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**
- ▶ Good heuristic design requires **in-depth knowledge** of the tactics and strategy of the game
- ▶ What if we don't possess such knowledge?

**Monte Carlo evaluation**

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**
- ▶ Generally designed to **converge in the limit**
  - ▶ An **infinite** number of samples would give an **exact** answer
  - ▶ As the **number of samples** increases, the **accuracy** of the answer improves
- ▶ Applications in physics, engineering, finance, weather forecasting, graphics, ...



# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
  - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay
- ▶ What we actually have are **pseudo-random number generators (PRNGs)**
- ▶ A PRNG is an algorithm which gives an **unpredictable** sequence of numbers based on a **seed**
- ▶ Sequence is **uniformly distributed**, i.e. all numbers have equal probability
- ▶ Seed is generally based on some source of **entropy**, e.g. system clock, mouse input, electronic noise

# Monte Carlo evaluation in games


- ▶ Based on **random rollouts**
  - while**  $s$  is not terminal **do**
    - let  $m$  be a random legal move from  $s$
    - update  $s$  by playing  $m$
- ▶ The **value** of a rollout is the **value** of the terminal state it reaches (i.e. 1 for a win,  $-1$  for a loss, 0 for a draw)
- ▶ Averaging gives the **expected value** of the initial state
- ▶ Higher expected value = more chance of winning

# Monte Carlo search


- ▶ **Flat Monte Carlo search:** 1-ply search with Monte Carlo evaluation
- ▶ How about minimax with  $d > 1$  and Monte Carlo evaluation?
  - ▶ Minimax assumes the evaluation is **deterministic**, but Monte Carlo is not
  - ▶ Not commonly used, mainly because there's something better...

# Monte Carlo Tree Search (MCTS)

# The perfect 1-ply search



# The multi-armed bandit problem



At each step pull one arm

Noisy/random reward signal

In order to:

- Minimise regret

- Maximise expected return  
(Find the best arm)

# Upper Confidence Bound (UCB1)

- Balance **exploitation** with **exploration**
- Select the arm that maximises

The diagram illustrates the UCB1 formula: 
$$\underbrace{\frac{V_a}{n_a}}_{\text{Exploitation}} + \underbrace{k \sqrt{\frac{\log n}{n_a}}}_{\text{Exploration}}$$
 Annotations include:


- An arrow from "Total reward from this arm so far" pointing to  $V_a$ .
- An arrow from "Total number of trials so far" pointing to  $\log n$ .
- An arrow from "Number of times this arm has been selected so far" pointing to  $n_a$  in the denominator of the square root term.
- An arrow from "Tuning constant" pointing to  $k$ .
- Brackets below the formula label the first term as "Exploitation" and the second term as "Exploration".

# UCB1 demo

<http://orangehelicopter.com/academic/bandits.html?ucb>




# Monte Carlo Tree Search (MCTS)



- Iteratively build a partial search tree, one node per iteration
- Monte Carlo evaluation (random playouts) for nonterminal states
- **Asymmetric**
  - balance exploitation with exploration
- **Anytime**
  - can do as many (or as few) iterations as we want
- **Aheuristic**
  - Monte Carlo evaluation requires only a forward model


# Monte Carlo Tree Search (MCTS)

Selection




Choose a node  
with unexpanded  
children

# Game Decisions



# Monte Carlo Tree Search (MCTS)



# Upper Confidence bound for Trees (UCT)

- Tree policy: UCB1  $\frac{V_a}{n_a} + k \sqrt{\frac{\log n}{n_a}}$
- Default policy: uniform random

# Demo



# Conclusion

- MCTS is a powerful **general-purpose** AI technique
  - Asymmetric, Anytime, Aheuristic
- MCTS has proven successful in several **challenging classes of games**
  - Games of imperfect information
  - Commercial mobile games
  - Real-time games
- It shows promise in **many other games and non-game applications**