



COMP250: Artificial Intelligence  
**5: Game Tree Search**

# Minimax search



# Minimax

# Minimax

- ▶ Terminal game states have a **value**

# Minimax

- ▶ Terminal game states have a **value**
  - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw

# Minimax

- ▶ Terminal game states have a **value**
  - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value

# Minimax

- ▶ Terminal game states have a **value**
  - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value

# Minimax

- ▶ Terminal game states have a **value**
  - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value
- ▶ Therefore I want to **maximise** the **minimum** value my opponent can achieve

# Minimax

- ▶ Terminal game states have a **value**
  - ▶ E.g. +1 for a win, -1 for a loss, 0 for a draw
- ▶ I want to **maximise** the value
- ▶ My opponent wants to **minimise** the value
- ▶ Therefore I want to **maximise** the **minimum** value my opponent can achieve
- ▶ This is generally only true for **two-player zero-sum** games

# Minimax search

# Minimax search

- ▶ Recursively defines a **value** for non-terminal game states

# Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move

# Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move
- ▶ If it’s my turn, the value is the **maximum** value over next states

# Minimax search

- ▶ Recursively defines a **value** for non-terminal game states
- ▶ Consider each possible “next state”, i.e. each possible move
- ▶ If it’s my turn, the value is the **maximum** value over next states
- ▶ If it’s my opponent’s turn, the value is the **minimum** value over next states

# Minimax search – example

minimax.pptx

# Minimax search pseudocode

```
procedure MINIMAX(state)
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
        for each possible nextState do
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
        for each possible nextState do
            v = MINIMAX(nextState)
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
        for each possible nextState do
            v = MINIMAX(nextState)
            bestValue = Max(bestValue, v)
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
        for each possible nextState do
            v = MINIMAX(nextState)
            bestValue = MAX(bestValue, v)
        return bestValue
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
        for each possible nextState do
            v = MINIMAX(nextState)
            bestValue = Max(bestValue, v)
        return bestValue
    else if state.currentPlayer = 2 then
```

# Minimax search pseudocode

```
procedure MINIMAX(state)
    if state is terminal then
        return value of state
    else if state.currentPlayer = 1 then
        bestValue = -∞
        for each possible nextState do
            v = MINIMAX(nextState)
            bestValue = MAX(bestValue, v)
        return bestValue
    else if state.currentPlayer = 2 then
        bestValue = +∞
        for each possible nextState do
            v = MINIMAX(nextState)
            bestValue = MIN(bestValue, v)
        return bestValue
```

# Stopping early

```
for each possible nextState do
    v = MINIMAX(nextState)
    bestValue = Max(bestValue, v)
```

# Stopping early

```
for each possible nextState do
    v = MINIMAX(nextState)
    bestValue = MAX(bestValue, v)
```

- State values are always between –1 and +1

# Stopping early

**for each** possible nextState **do**

$v = \text{MINIMAX}(\text{nextState})$

$\text{bestValue} = \text{MAX}(\text{bestValue}, v)$

- ▶ State values are always between  $-1$  and  $+1$
- ▶ So if we ever have  $\text{bestValue} = 1$ , we can stop early

# Stopping early

**for each** possible nextState **do**

$v = \text{MINIMAX}(\text{nextState})$

$\text{bestValue} = \text{MAX}(\text{bestValue}, v)$

- ▶ State values are always between  $-1$  and  $+1$
- ▶ So if we ever have  $\text{bestValue} = 1$ , we can stop early
- ▶ Similarly when minimising if  $\text{bestValue} = -1$

# Using minimax search

# Using minimax search

- To decide what move to play next...

# Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move

# Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move
- ▶ Choose the move with the maximum score

# Using minimax search

- ▶ To decide what move to play next...
- ▶ Calculate the minimax value for each move
- ▶ Choose the move with the maximum score
- ▶ If there are several with the same score, choose one at random

# Minimax and game theory

# Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**

# Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**

# Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**
- ▶ I.e. a minimax player plays **perfectly**

# Minimax and game theory

- ▶ For a **two-player zero-sum** game with **perfect information** and **sequential moves**
- ▶ Minimax search will always find a **Nash equilibrium**
- ▶ I.e. a minimax player plays **perfectly**
- ▶ **But...**

# Minimax for larger games

# Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states

# Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully

# Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully
  - ▶ Connect 4 has  $\approx 10^{13}$  states

# Minimax for larger games

- ▶ The game tree for noughts and crosses has only a few thousand states
- ▶ Most games are too large to search fully
  - ▶ Connect 4 has  $\approx 10^{13}$  states
  - ▶ Chess has  $\approx 10^{47}$  states

# Heuristics for search



# Depth limiting

# Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states

# Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games

# Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as +1 / 0 / -1

# Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as +1 / 0 / -1
- ▶ For nonterminal states at depth  $d$ , apply a heuristic evaluation instead of searching deeper

# Depth limiting

- ▶ Standard minimax needs to search all the way to **terminal** (game over) states
- ▶ **Depth limiting** is a common technique to apply minimax to larger games
- ▶ Still evaluate terminal states as +1 / 0 / -1
- ▶ For nonterminal states at depth  $d$ , apply a heuristic evaluation instead of searching deeper
- ▶ Evaluation is a number between -1 and +1, estimating the probable outcome of the game

# 1-ply search

# 1-ply search

- Case  $d = 1$

# 1-ply search

- ▶ Case  $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move

# 1-ply search

- ▶ Case  $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move
- ▶ This is computationally fast

# 1-ply search

- ▶ Case  $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move
- ▶ This is computationally fast
- ▶ Often easier to design a “which state is better” heuristic than to directly design a “which move to play” heuristic

# 1-ply search

- ▶ Case  $d = 1$
- ▶ For each move, evaluate the state resulting from playing that move
- ▶ This is computationally fast
- ▶ Often easier to design a “which state is better” heuristic than to directly design a “which move to play” heuristic
- ▶ This is essentially a **utility-based AI**

# Move ordering

# Move ordering

- Minimax can **stop early** if it sees a value of +1 for maximising player or -1 for minimising player

# Move ordering

- ▶ Minimax can **stop early** if it sees a value of +1 for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this

# Move ordering

- ▶ Minimax can **stop early** if it sees a value of +1 for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search

# Move ordering

- ▶ Minimax can **stop early** if it sees a value of +1 for maximising player or -1 for minimising player
- ▶ Modifications to minimax algorithm (e.g. **alpha-beta pruning**) lead to more of this
- ▶ Thus ordering moves from **best to worst** means faster search
- ▶ How do we know which moves are “best” and “worst”? Use a heuristic!

# Designing heuristics

# Designing heuristics

- The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**

# Designing heuristics

- ▶ The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**
- ▶ Good heuristic design requires **in-depth knowledge** of the tactics and strategy of the game

# Designing heuristics

- ▶ The **playing strength** of depth limited minimax depends heavily on the design of the **heuristic**
- ▶ Good heuristic design requires **in-depth knowledge** of the tactics and strategy of the game
- ▶ What if we don't possess such knowledge?

# Monte Carlo evaluation



# Monte Carlo methods

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**
- ▶ Generally designed to **converge in the limit**

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**
- ▶ Generally designed to **converge in the limit**
  - ▶ An **infinite** number of samples would give an **exact** answer

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**
- ▶ Generally designed to **converge in the limit**
  - ▶ An **infinite** number of samples would give an **exact** answer
  - ▶ As the **number of samples** increases, the **accuracy** of the answer improves

# Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**
- ▶ Generally designed to **converge in the limit**
  - ▶ An **infinite** number of samples would give an **exact** answer
  - ▶ As the **number of samples** increases, the **accuracy** of the answer improves
- ▶ Applications in physics, engineering, finance, weather forecasting, graphics, ...

# Aside: “randomness” in computing

# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness

# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
  - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay

# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
  - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay
- ▶ What we actually have are **pseudo-random number generators (PRNGs)**

# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
  - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay
- ▶ What we actually have are **pseudo-random number generators (PRNGs)**
- ▶ A PRNG is an algorithm which gives an **unpredictable** sequence of numbers based on a **seed**

# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
  - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay
- ▶ What we actually have are **pseudo-random number generators (PRNGs)**
- ▶ A PRNG is an algorithm which gives an **unpredictable** sequence of numbers based on a **seed**
- ▶ Sequence is **uniformly distributed**, i.e. all numbers have equal probability

# Aside: “randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
  - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay
- ▶ What we actually have are **pseudo-random number generators (PRNGs)**
- ▶ A PRNG is an algorithm which gives an **unpredictable** sequence of numbers based on a **seed**
- ▶ Sequence is **uniformly distributed**, i.e. all numbers have equal probability
- ▶ Seed is generally based on some source of **entropy**, e.g. system clock, mouse input, electronic noise

# Monte Carlo evaluation in games

# Monte Carlo evaluation in games

- ▶ Based on **random rollouts**

# Monte Carlo evaluation in games

- Based on **random rollouts**

**while**  $s$  is not terminal **do**

    let  $m$  be a random legal move from  $s$

    update  $s$  by playing  $m$

# Monte Carlo evaluation in games

- ▶ Based on **random rollouts**

**while**  $s$  is not terminal **do**

    let  $m$  be a random legal move from  $s$

    update  $s$  by playing  $m$

- ▶ The **value** of a rollout is the **value** of the terminal state it reaches (i.e. 1 for a win, -1 for a loss, 0 for a draw)

# Monte Carlo evaluation in games

- ▶ Based on **random rollouts**

**while**  $s$  is not terminal **do**

    let  $m$  be a random legal move from  $s$

    update  $s$  by playing  $m$

- ▶ The **value** of a rollout is the **value** of the terminal state it reaches (i.e. 1 for a win, -1 for a loss, 0 for a draw)
- ▶ Averaging gives the **expected value** of the initial state

# Monte Carlo evaluation in games

- ▶ Based on **random rollouts**

**while**  $s$  is not terminal **do**

    let  $m$  be a random legal move from  $s$

    update  $s$  by playing  $m$

- ▶ The **value** of a rollout is the **value** of the terminal state it reaches (i.e. 1 for a win, -1 for a loss, 0 for a draw)
- ▶ Averaging gives the **expected value** of the initial state
- ▶ Higher expected value = more chance of winning

# Monte Carlo search

# Monte Carlo search

- ▶ **Flat Monte Carlo search:** 1-ply search with Monte Carlo evaluation

# Monte Carlo search

- ▶ **Flat Monte Carlo search:** 1-ply search with Monte Carlo evaluation
- ▶ How about minimax with  $d > 1$  and Monte Carlo evaluation?

# Monte Carlo search

- ▶ **Flat Monte Carlo search:** 1-ply search with Monte Carlo evaluation
- ▶ How about minimax with  $d > 1$  and Monte Carlo evaluation?
  - ▶ Minimax assumes the evaluation is **deterministic**, but Monte Carlo is not

# Monte Carlo search

- ▶ **Flat Monte Carlo search:** 1-ply search with Monte Carlo evaluation
- ▶ How about minimax with  $d > 1$  and Monte Carlo evaluation?
  - ▶ Minimax assumes the evaluation is **deterministic**, but Monte Carlo is not
  - ▶ Not commonly used, mainly because there's something better...