

COMP250: Artificial Intelligence

5: Monte Carlo Tree Search

Assignment check-in

AI component

- ▶ **Peer review** will run during **development week** (next week)
- ▶ Submit your component for review by **11:59PM Monday 5th March**
- ▶ Complete your reviews by **11:59PM Friday 9th March**

Research wiki

MicroRTS bot

- ▶ The server is now live!
- ▶ See the `comp250-bot` repository on GitHub for details

Monte Carlo evaluation

From last time

- ▶ It is useful to have a **heuristic evaluation function** for nonterminal states
- ▶ Allows 1-ply search, depth-limited minimax, ...
- ▶ Designing a good heuristic requires in-depth knowledge of the game
- ▶ What if you don't have such knowledge?

Expected value

- ▶ Let X be a **random variable**
- ▶ Let $p(x)$ be the probability that X has value x
- ▶ Then the **expected value** of X is

$$\sum_x x \cdot p(x)$$

Expected value — example

- ▶ A slot machine pays out:
 - ▶ £1 with probability 0.05
 - ▶ £5 with probability 0.03
 - ▶ £10 with probability 0.02
 - ▶ Nothing with probability 0.9
- ▶ The expected payout is

$$1 \times 0.05 + 5 \times 0.03 + 10 \times 0.02 + 0 \times 0.9 = 0.4$$

i.e. £0.40

- ▶ What this means: if you play the slot machine N times, on average you will win $N \times £0.40$

“Randomness” in computing

- ▶ Digital computers are **deterministic**, so there's no such thing as true randomness
 - ▶ Cryptographically secure systems use an external source of randomness e.g. atmospheric noise, radioactive decay
- ▶ What we actually have are **pseudo-random number generators (PRNGs)**
- ▶ A PRNG is an algorithm which gives an **unpredictable** sequence of numbers based on a **seed**
- ▶ Sequence is **uniformly distributed**, i.e. all numbers have equal probability
- ▶ Seed is generally based on some source of **entropy**, e.g. system clock, mouse input, electronic noise

Monte Carlo methods

- ▶ In computing, a **Monte Carlo method** is an algorithm based on **averaging over random samples**
- ▶ The **average** over a large number of samples is a good approximation of the **expected value**
- ▶ Used for **quickly approximating** quantities over **large domains**
- ▶ Generally designed to **converge in the limit**
 - ▶ An **infinite** number of samples would give an **exact** answer
 - ▶ As the **number of samples** increases, the **accuracy** of the answer improves
- ▶ Applications in physics, engineering, finance, weather forecasting, graphics, ...

Monte Carlo evaluation in games

- ▶ Based on **random rollouts**

while s is not terminal **do**

 let m be a random legal move from s

 update s by playing m

- ▶ The **value** of a rollout is the **value** of the terminal state it reaches (i.e. 1 for a win, -1 for a loss, 0 for a draw)
- ▶ Averaging gives the **expected value** of the initial state
- ▶ Higher expected value = more chance of winning

Monte Carlo search

- ▶ **Flat Monte Carlo search:** 1-ply search with Monte Carlo evaluation
- ▶ How about minimax with $d > 1$ and Monte Carlo evaluation?
 - ▶ Minimax assumes the evaluation is **deterministic**, but Monte Carlo is not
 - ▶ Not commonly used, mainly because there's something better...

Monte Carlo Tree Search


Monte Carlo Tree Search (MCTS)

- ▶ Like Monte Carlo evaluation, based on **rollouts**
- ▶ First few rollouts are **random**
- ▶ However, statistics from these rollouts are used to **bias** future rollouts
- ▶ Bias rollouts towards **plausible** lines of play, i.e. where each player is trying to play the best move

The MCTS algorithm

- ▶ MCTS builds a **tree**
- ▶ Initially, the tree consists of a single **root node**
- ▶ Each rollout has four stages:
 - ▶ **Selection**: Starting from the root, descend the tree by choosing moves. Continue until we reach a node which does not yet have children for all legal moves.
 - ▶ **Expansion**: Choose a random legal move for which the current node does not have a child node. Add this new node to the tree.
 - ▶ **Simulation**: Perform a Monte Carlo rollout, playing random moves until a terminal state is reached.
 - ▶ **Backpropagation**: For each node visited during **selection** and **expansion**, update the node's statistics based on the result of the simulation.
- ▶ Perform many rollouts, then use the statistics at the top level of the tree to choose the best move

The MCTS algorithm



Selection policy

- ▶ Selection must balance:
 - ▶ **Exploitation** of moves that are known to be good
 - ▶ **Exploration** of moves that have not often been tried
- ▶ This can be modelled as a **multi-armed bandit problem**

Multi-armed bandits

- ▶ We have a row of one-armed bandits (slot machines)
- ▶ We **do not know** the payout probabilities of any of them, and they're all different
- ▶ How to maximise our winnings?
- ▶ Again must balance
 - ▶ **Exploitation** of machines that are known to have a high expected payout
 - ▶ **Exploration** of machines that have not been tried often, to get a better estimate of their expected payout

Upper Confidence Bound (UCB)

- ▶ For each machine m , record:
 - ▶ n_m : the number of plays of this machine
 - ▶ V_m : the total winnings from playing this machine
 - ▶ $n = \sum_m n_m$, total number of plays across all machines
- ▶ At each stage, play the machine for which

$$\frac{V_m}{n_m} + c\sqrt{\frac{\log n}{n_m}}$$

is largest

- ▶ $\frac{V_m}{n_m}$ is the **exploitation** part: average payout from this machine so far
- ▶ $\sqrt{\frac{\log n}{n_m}}$ is the **exploration** part: large if n_m is small
- ▶ c is a parameter for adjusting the balance between exploitation and exploration

UCB demo

[http://orangehelicopter.com/academic/bandits.
html?ucb](http://orangehelicopter.com/academic/bandits.html?ucb)

Upper Confidence Bound for Trees (UCT)

- ▶ Use UCB as the selection policy
- ▶ In each node x , record:
 - ▶ n_x : the number of visits to this node
 - ▶ V_x : the total value of rollouts through this node
- ▶ From node p , choose the child q such that

$$\frac{V_q}{n_q} + c \sqrt{\frac{\log n_p}{n_q}}$$

is largest

UCT demo

Benefits of MCTS

- ▶ “Vanilla” MCTS is **game independent**
- ▶ But if game-specific heuristics are available, they can be used to **enhance** MCTS
- ▶ MCTS is **anytime**
 - ▶ Can stop it after **any** amount of computation (within reason) and get a reasonably good answer
 - ▶ Compare with minimax: $O(e^d)$ for depth d
- ▶ Does not suffer from **horizon effect**
 - ▶ Minimax at depth d cannot “see” what happens $d + 1$ moves in the future
 - ▶ MCTS can build the tree as deep as it likes
 - ▶ Selects which parts of the tree to expand more deeply

MCTS for games of imperfect information