

COMP250: Artificial Intelligence


## **9: Deep learning**

**Neural networks**

# Artificial Neural Networks (ANNs)


- ▶ **Inspired by** the structure of biological brains
- ▶ Idea has been around since the 1950s
- ▶ Recent resurgence of interest: today's powerful CPUs and GPUs allow much larger ANNs to be used

# Real neurons



- ▶ An **electrically excitable** cell
  - ▶ Neurons are **connected together**
  - ▶ Connections can be **excitatory** or **inhibitory**
- 
- ▶ If enough excitatory signals are received, the neuron **fires** — sends an electrical signal to the connected neurons
  - ▶ Human brain contains approximately **100 billion** neurons

# An artificial neuron



- ▶ A **perceptron**
- ▶ Inputs  $x_1, \dots, x_m$  are outputs from **other perceptrons**
- ▶ Each input has a **weight**  $w_i$  between  $-1$  and  $+1$

# Perceptron activation


- ▶ The perceptron calculates a **weighted sum**

$$w_0 + w_1x_1 + \cdots + w_mx_m$$

- ▶ This goes through an **activation function**
- ▶ Simplest: **step function**


$$\text{output} = \begin{cases} 1 & \text{if sum} \geq \text{threshold} \\ 0 & \text{if sum} < \text{threshold} \end{cases}$$

- ▶ More common: **sigmoid function**



# An artificial neural network

Input layer      Hidden Layers      Output Layer



- ▶ A **multilayer perceptron (MLP)**
- ▶ Consists of an **input layer**, several **hidden layers** and an **output layer**
- ▶ Each layer is an array of **perceptrons**
- ▶ Each perceptron's output is connected to **every** perceptron in the next layer

# Image classification



- ▶ Classic example:  
**handwritten digit recognition**
- ▶ Given a **raster image**, which of the digits 0 to 9 does it represent?





<https://twitter.com/NaughtThought/status/846262063827730432>

# MLPs for image classification

- ▶ **Input:** pixels of the image, reduced down to 1 bit per pixel (i.e. black or white)
  - ▶ Input layer: 1 perceptron per pixel
- ▶ **Output:** 10 bits corresponding to digits 0 to 9, of which exactly one should be set
  - ▶ Output layer: 10 perceptrons
- ▶ **Hidden layers:** ???
  - ▶ Parameters to tune
- ▶ **Weights:** ???

# How to set the weights?

- ▶ We need to **train** the network
- ▶ Idea:
  - ▶ Feed in **training data**
  - ▶ When the network happens to give the correct answer, **reinforce** the relevant weights
  - ▶ Repeat until a desired **accuracy** is obtained
- ▶ Note: this requires a large amount of training data that is **tagged**, i.e. for which we already know the correct answer

# Stochastic gradient descent

- ▶ **Gradient descent**: opposite of **gradient ascent** a.k.a. **hillclimbing**
- ▶ Want to minimise the **error** over the training data
- ▶ **Stochastic**: perform several training **epochs**
- ▶ Each epoch uses a randomly sampled **subset** of the training data
- ▶ This reduces computation time, and helps to escape local optima

# ANN example

<http://playground.tensorflow.org>

# Overfitting

- ▶ ANN learns **patterns** in the training data
- ▶ Insufficient training data might result in the network learning “patterns” that are actually random anomalies


**Deep learning**

# Deep learning

- ▶ Basically, the use of large ANNs with **many layers**
- ▶ Often uses **large training sets**
- ▶ Training often uses powerful **GPUs** — many times faster than training on the CPU





# Convolutional Neural Networks (ConvNets)



- Layers are **2D arrays**
- Neurons in convolutional layers are only connected to nearby neurons
- There are also fully connected layers


Deep neural networks learn hierarchical feature representations



# DeepDream

- ▶ Train a ConvNet to recognise something (e.g. faces, objects, animals)
- ▶ Run the network in “reverse”
  - ▶ Adjust the image (e.g. via gradient ascent) so that it is more strongly recognised by the network

# DeepDream



# Style transfer

- ▶ Train a ConvNet to recognise a particular artistic style
- ▶ Run the network in “reverse” on an input image
  - ▶ Adjust the image (e.g. via gradient ascent) so that it is more strongly recognised by the network

## Style transfer



Source image (**Style**)



Target image (**Content**)




Output ([deepart](#))

A Neural Algorithm of Artistic Style [[Gatys et al. 2015](#)]

# Generative Adversarial Networks (GANs)

- ▶ Two ANNs trained in parallel
  - ▶ One to generate “fake” artefacts
  - ▶ One to distinguish “real” from “fake”
- ▶ [http://research.nvidia.com/publication/2017-10\\_Progressive-Growing-of](http://research.nvidia.com/publication/2017-10_Progressive-Growing-of)

# Learning to play Atari games (Mnih et al, 2015)





# AlphaGo (Silver et al, 2017)

- ▶ MCTS with ANNs for move pruning, simulation playouts and state evaluation
- ▶ ANNs trained on both expert human matches and self-play (reinforcement learning)
- ▶ Defeated Lee Sedol, world Go champion

# AlphaZero (Silver et al, 2018)

- ▶ Similar MCTS+ANN architecture to AlphaGo
- ▶ Trained by reinforcement learning (self-play) only
- ▶ After only 9 hours\* of training, defeated Stockfish (one of the strongest chess programs available) in a 100-match tournament
  - ▶ \* On a cluster of 5000 of Google's custom Tensor Processing Units
- ▶ Stockfish is based on decades of research by expert chess players and AI programmers — AlphaZero started from no chess-specific knowledge whatsoever (other than the rules of the game)

# **Deep learning for PCG**

<https://www.youtube.com/watch?v=3wcpLwvBTYo>