

COMP350: Algorithms & Optimisation

4: GPU Optimisation

Learning outcomes


By the end of today's session, you will be able to:

- ▶ **Recall** the key stages of the graphics pipeline
- ▶ **Understand** the GPU Debugger
- ▶ **Explain** some of the key areas for optimisation


The 3D graphics pipeline



The 3D graphics pipeline




Vertex processing




Vertex processing

- ▶ Geometry is provided to the GPU as a **mesh** of **triangles**




Vertex processing




- ▶ Geometry is provided to the GPU as a **mesh** of **triangles**
- ▶ Each triangle has three **vertices** specified in 3D space (x, y, z)

Vertex processing




- ▶ Geometry is provided to the GPU as a **mesh** of **triangles**
- ▶ Each triangle has three **vertices** specified in 3D space (x, y, z)
- ▶ Vertex processor **transforms** (rotates, moves, scales) vertices and **projects** them into 2D screen space (x, y)

Vertex processing




- ▶ Geometry is provided to the GPU as a **mesh** of **triangles**
- ▶ Each triangle has three **vertices** specified in 3D space (x, y, z)
- ▶ Vertex processor **transforms** (rotates, moves, scales) vertices and **projects** them into 2D screen space (x, y)
- ▶ May also apply particle simulations, skeletal animations or deformations, etc.

Rasterisation




Rasterisation




- Determine **which fragments** are covered by the triangle

Rasterisation




- ▶ Determine **which fragments** are covered by the triangle
- ▶ In practical terms, “fragment” = “pixel”

Rasterisation




- ▶ Determine **which fragments** are covered by the triangle
- ▶ In practical terms, “fragment” = “pixel”
- ▶ Vertex processor can associate **data** with each vertex; this is **interpolated** across the fragments

Fragment processing




Fragment processing




- Determine the **colour** of each fragment covered by the triangle

Fragment processing




- Determine the **colour** of each fragment covered by the triangle
- **Textures** are 2D images that can be **wrapped** onto a 3D object

Fragment processing




- Determine the **colour** of each fragment covered by the triangle
- **Textures** are 2D images that can be **wrapped** onto a 3D object
- Colour is calculated based on **texture**, **lighting** and other properties of the surface being rendered (e.g. shininess, roughness)

Blending




Blending

- Combine these fragments with the existing content of the image buffer




Blending



- ▶ Combine these fragments with the existing content of the image buffer
- ▶ **Depth testing:** if the new fragment is "in front" of the old one, replace it; if it is "behind", discard it

Blending



- ▶ Combine these fragments with the existing content of the image buffer
- ▶ **Depth testing:** if the new fragment is "in front" of the old one, replace it; if it is "behind", discard it
- ▶ **Alpha blending:** combine the old and new colours for a semi-transparent appearance

Standard Shaders

Standard Shaders

- ▶ The vertex processor and fragment processor are **programmable**

Standard Shaders

- ▶ The vertex processor and fragment processor are **programmable**
- ▶ Programs for these units are called **shaders**

Standard Shaders

- ▶ The vertex processor and fragment processor are **programmable**
- ▶ Programs for these units are called **shaders**
- ▶ **Vertex shader**: responsible for geometric transformations, deformations, and projection

Standard Shaders

- ▶ The vertex processor and fragment processor are **programmable**
- ▶ Programs for these units are called **shaders**
- ▶ **Vertex shader**: responsible for geometric transformations, deformations, and projection
- ▶ **Fragment shader**: responsible for the visual appearance of the surface

Standard Shaders

- ▶ The vertex processor and fragment processor are **programmable**
- ▶ Programs for these units are called **shaders**
- ▶ **Vertex shader**: responsible for geometric transformations, deformations, and projection
- ▶ **Fragment shader**: responsible for the visual appearance of the surface
- ▶ Vertex shader and fragment shader are separate programs, but the vertex shader can pass arbitrary values through to the fragment shader

Other Shaders

Other Shaders

► Geometry Shader:

Other Shaders

- ▶ **Geometry Shader:**

- ▶ Operates on primitives

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives
- Usually used to expand geometry (i.e take in one point and produce a triangle)

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives
- Usually used to expand geometry (i.e take in one point and produce a triangle)
- Typically used for fur, hair, particle systems

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives
- Usually used to expand geometry (i.e take in one point and produce a triangle)
- Typically used for fur, hair, particle systems

► **Tessellation Control Shader:**

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives
- Usually used to expand geometry (i.e take in one point and produce a triangle)
- Typically used for fur, hair, particle systems

► **Tessellation Control Shader:**

- Receives input from vertex shader

Other Shaders

▶ **Geometry Shader:**

- ▶ Operates on primitives
- ▶ Can emit zero, one or more primitives
- ▶ Usually used to expand geometry (i.e take in one point and produce a triangle)
- ▶ Typically used for fur, hair, particle systems

▶ **Tessellation Control Shader:**

- ▶ Receives input from vertex shader
- ▶ Determines the amount of tessellation on a primitive

Other Shaders

▶ **Geometry Shader:**

- ▶ Operates on primitives
- ▶ Can emit zero, one or more primitives
- ▶ Usually used to expand geometry (i.e take in one point and produce a triangle)
- ▶ Typically used for fur, hair, particle systems

▶ **Tessellation Control Shader:**

- ▶ Receives input from vertex shader
- ▶ Determines the amount of tessellation on a primitive
- ▶ Perform any transformation on the patch data

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives
- Usually used to expand geometry (i.e take in one point and produce a triangle)
- Typically used for fur, hair, particle systems

► **Tessellation Control Shader:**

- Receives input from vertex shader
- Determines the amount of tessellation on a primitive
- Perform any transformation on the patch data
- Can change the size of the patch, add more vertices or fewer

Other Shaders

► **Geometry Shader:**

- Operates on primitives
- Can emit zero, one or more primitives
- Usually used to expand geometry (i.e take in one point and produce a triangle)
- Typically used for fur, hair, particle systems

► **Tessellation Control Shader:**

- Receives input from vertex shader
- Determines the amount of tessellation on a primitive
- Perform any transformation on the patch data
- Can change the size of the patch, add more vertices or fewer
- Typically used for level of detail and stitching

Other Shaders

Other Shaders

► Tessellation Evaluation Shader:

Other Shaders

- ▶ **Tessellation Evaluation Shader:**
 - ▶ Relieves input from the Tessellator

Other Shaders

► Tessellation Evaluation Shader:

- Relieves input from the Tessellator
- Calculates the new vertices in the patch

Other Shaders

▶ Tessellation Evaluation Shader:

- ▶ Relieves input from the Tessellator
- ▶ Calculates the new vertices in the patch
- ▶ Works in conjunction with the TC Shaders

Other Shaders

- ▶ **Tessellation Evaluation Shader:**
 - ▶ Relieves input from the Tessellator
 - ▶ Calculates the new vertices in the patch
 - ▶ Works in conjunction with the TC Shaders
- ▶ **Compute Shader:**

Other Shaders

▶ **Tessellation Evaluation Shader:**

- ▶ Relieves input from the Tessellator
- ▶ Calculates the new vertices in the patch
- ▶ Works in conjunction with the TC Shaders

▶ **Compute Shader:**

- ▶ These types of shaders allow you to carry out general purpose computing on the GPU

Other Shaders

▶ **Tessellation Evaluation Shader:**

- ▶ Relieves input from the Tessellator
- ▶ Calculates the new vertices in the patch
- ▶ Works in conjunction with the TC Shaders

▶ **Compute Shader:**

- ▶ These types of shaders allow you to carry out general purpose computing on the GPU
- ▶ Can access all the same data as normal shaders, exceptions are attributes

Other Shaders

► Tessellation Evaluation Shader:

- Relieves input from the Tessellator
- Calculates the new vertices in the patch
- Works in conjunction with the TC Shaders

► Compute Shader:

- These types of shaders allow you to carry out general purpose computing on the GPU
- Can access all the same data as normal shaders, exceptions are attributes
- The shader has to write to an image or a shader storage object

Other Shaders

► Tessellation Evaluation Shader:

- Relieves input from the Tessellator
- Calculates the new vertices in the patch
- Works in conjunction with the TC Shaders

► Compute Shader:

- These types of shaders allow you to carry out general purpose computing on the GPU
- Can access all the same data as normal shaders, exceptions are attributes
- The shader has to write to an image or a shader storage object
- This shader type can do simulations, AI or any other general purpose processing

GPU Profiling



Live Demo

Live Demo

► Render Doc

Live Demo

- ▶ Render Doc
- ▶ Unity

Live Demo

- ▶ Render Doc
- ▶ Unity
- ▶ Unreal

GPU Optimisation



Visibility Culling

Visibility Culling

- ▶ You should always cull your scene based on the cameras view fulstrum

Visibility Culling

- ▶ You should always cull your scene based on the cameras view frustum
- ▶ This will allow to eliminate objects that are not visible

Visibility Culling

- ▶ You should always cull your scene based on the cameras view frustum
- ▶ This will allow to eliminate objects that are not visible
- ▶ This combined with a scene graph will allow us to cull large parts of the scene

Visibility Culling

- ▶ You should always cull your scene based on the cameras view frustum
- ▶ This will allow to eliminate objects that are not visible
- ▶ This combined with a scene graph will allow us to cull large parts of the scene
- ▶ You should also sort all visible objects from back to front

Visibility Culling

- ▶ You should always cull your scene based on the cameras view frustum
- ▶ This will allow to eliminate objects that are not visible
- ▶ This combined with a scene graph will allow us to cull large parts of the scene
- ▶ You should also sort all visible objects from back to front
- ▶ Caveat, transparent object should be sorted front to back

State Changes

State Changes

- ▶ You should attempt to minimize state changes

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes
 - ▶ Changing Shaders/Materials

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes
 - ▶ Changing Shaders/Materials
 - ▶ Changing Pipeline States

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes
 - ▶ Changing Shaders/Materials
 - ▶ Changing Pipeline States
 - ▶ Changing Active Textures

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes
 - ▶ Changing Shaders/Materials
 - ▶ Changing Pipeline States
 - ▶ Changing Active Textures
- ▶ If you are working in an engine, try to minimum the amount of different materials

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes
 - ▶ Changing Shaders/Materials
 - ▶ Changing Pipeline States
 - ▶ Changing Active Textures
- ▶ If you are working in an engine, try to minimum the amount of different materials
- ▶ This will allow the engine to sort the render queue based on material

State Changes

- ▶ You should attempt to minimize state changes
- ▶ This includes
 - ▶ Changing Shaders/Materials
 - ▶ Changing Pipeline States
 - ▶ Changing Active Textures
- ▶ If you are working in an engine, try to minimum the amount of different materials
- ▶ This will allow the engine to sort the render queue based on material
- ▶ Attempt to use a texture atlas to manage your textures

Batching

Batching

- ▶ You should attempt to batch your geometry to minimise draw calls

Batching

- ▶ You should attempt to batch your geometry to minimise draw calls
 - ▶ Unity: Mark GameObject as static

Batching

- ▶ You should attempt to batch your geometry to minimise draw calls
 - ▶ Unity: Mark GameObject as static
 - ▶ Unreal: Mobility settings, change actors to Static

Batching

- ▶ You should attempt to batch your geometry to minimise draw calls
 - ▶ Unity: Mark GameObject as static
 - ▶ Unreal: Mobility settings, change actors to Static
 - ▶ OpenGL: Have only a few larger VBOs

Instancing

Instancing

- ▶ This allows the GPU to draw one version of the mesh multiple times in one draw call

Instancing

- ▶ This allows the GPU to draw one version of the mesh multiple times in one draw call
 - ▶ Unity: based on materials - Enable Instancing

Instancing

- ▶ This allows the GPU to draw one version of the mesh multiple times in one draw call
 - ▶ Unity: based on materials - Enable Instancing
 - ▶ Unreal: <https://forums.unrealengine.com/unreal-engine/events/90518-canon-man-tutorial-hierarchical-instanced-118229>

Instancing

- ▶ This allows the GPU to draw one version of the mesh multiple times in one draw call
 - ▶ Unity: based on materials - Enable Instancing
 - ▶ Unreal: <https://forums.unrealengine.com/unreal-engine/events/90518-canon-man-tutorial-hierarchical-instanced-118229>
 - ▶ OpenGL: <http://www.opengl-tutorial.org/intermediate-tutorials/billboards-particles/particles-instancing/>

Shaders

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems
 - ▶ Reduce Texture reads

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems
 - ▶ Reduce Texture reads
 - ▶ Pack Texture data

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems
 - ▶ Reduce Texture reads
 - ▶ Pack Texture data
 - ▶ Reduce bandwidth by using a 16 bit texture

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems
 - ▶ Reduce Texture reads
 - ▶ Pack Texture data
 - ▶ Reduce bandwidth by using a 16 bit texture
- ▶ Avoid branching

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems
 - ▶ Reduce Texture reads
 - ▶ Pack Texture data
 - ▶ Reduce bandwidth by using a 16 bit texture
- ▶ Avoid branching
 - ▶ Prefer static branching over dynamic

Shaders

- ▶ Texture Sampling is one of the biggest source of memory access problems
 - ▶ Reduce Texture reads
 - ▶ Pack Texture data
 - ▶ Reduce bandwidth by using a 16 bit texture
- ▶ Avoid branching
 - ▶ Prefer static branching over dynamic
 - ▶ With static branching, the loop can be evaluated at compile time